NVIDIA OptiX 6.0
Programming Guide

8 January 2019
Version 6.0

NVIDIA OptiX 6.0 – Programming Guide

Copyright Information

Contents

OptiX 6.0 Programming Guide © 2019 NVIDIA Corporation

1   OptiX overview

GPUs are best at exploiting very high degrees of parallelism, and ray tracing fits that requirement perfectly. However, typical ray tracing algorithms can be highly irregular, which poses serious challenges for anyone trying to exploit the full raw computational potential of a GPU. The NVIDIA® OptiX™ ray tracing engine and API address those challenges and provide a framework for harnessing the enormous computational power of both current- and future-generation graphics hardware to incorporate ray tracing into interactive applications. By using OptiX together with NVIDIA® CUDA® architecture, interactive ray tracing is finally feasible for developers without a Ph.D. in computer graphics and a team of ray tracing engineers.

OptiX is not itself a renderer. Instead, it is a scalable framework for building ray tracing based applications. The OptiX engine is composed of two symbiotic parts: 1) a host-based API that defines data structures for ray tracing, and 2) a CUDA C++-based programming system that can produce new rays, intersect rays with surfaces, and respond to those intersections. Together, these two pieces provide low-level support for "raw ray tracing." This allows user-written applications that use ray tracing for graphics, collision detection, sound propagation, visibility determination, etc.

1.1   Motivation

By abstracting the execution model of a generic ray tracer, OptiX makes it easier to assemble a ray tracing system, leveraging custom-built algorithms for object traversal, shader dispatch and memory management. Furthermore, the resulting system will be able to take advantage of future evolution in GPU hardware and OptiX SDK releases—similar to the manner that OpenGL and Direct3D provide an abstraction for the rasterization pipeline.

Wherever possible, the OptiX engine avoids specification of ray tracing behaviors and instead provides mechanisms to execute user- provided CUDA C code to implement shading (including recursive rays), camera models, and even color representations. Consequently, the OptiX engine can be used for Whitted-style ray tracing, path tracing, collision detection, photon mapping, or any other ray tracing-based algorithm. It is designed to operate either standalone or in conjunction with an OpenGL or DirectX application for hybrid ray tracing-rasterization applications.

1.2   Programming model

At the core of OptiX is a simple but powerful abstract model of a ray tracer. This ray tracer employs user-provided programs to control the initiation of rays, intersection of rays with surfaces, shading with materials, and spawning of new rays. Rays carry user-specified payloads that describe per-ray variables such as color, recursion depth, importance, or other attributes. Developers provide these functions to OptiX in the form of CUDA C-based functions. Because ray tracing is an inherently recursive algorithm, OptiX allows user programs to recursively spawn new rays, and the internal execution mechanism manages all the details of a recursion stack. OptiX also provides flexible dynamic function dispatch and a

sophisticated variable inheritance mechanism so that ray tracing systems can be written very generically and compactly.

1.3   Ray tracing basics

"Ray tracing" is an overloaded term whose meaning can depend on context. Sometimes it refers to the computation of the intersection points between a 3D line and a set of 3D objects such as spheres. Sometimes it refers to a specific algorithm such as Whitted's method of generating pictures or the oil exploration industry's algorithm for simulating ground wave propagation. Other times it refers to a family of algorithms that include Whitted's algorithm along with others such as distribution ray tracing. OptiX is a ray tracing engine in the first sense of the word: it allows the user to intersect rays and 3D objects. As such it can be used to build programs that fit the other use of "ray tracing" such as Whitted's algorithm. In addition OptiX provides the ability for users to write their own programs to generate rays and to define behavior for when rays hit objects.

For graphics, ray tracing was originally proposed by Arthur Appel in 1968 for rendering solid objects. In 1980, Turner Whitted pursued the idea further by introducing recursion to enable reflective and refractive effects. Subsequent advances in ray tracing increased accuracy by introducing effects for depth of field, diffuse inter- reflection, soft shadows, motion blur, and other optical effects. Simultaneously, numerous researchers have improved the performance of ray tracing using new algorithms for indexing the objects in the scene.

Realistic rendering algorithms based on ray tracing have been used to accurately simulate light transport. Some of these algorithms simulate the propagation of photons in a virtual environment. Others follow adjoint photons "backward" from a virtual camera to determine where they originated. Still other algorithms use bidirectional methods. OptiX operates at a level below such algorithmic decisions, so can be used to build any of those algorithms.

Ray tracing has often been used for non-graphics applications. In the computer-aided design community, ray tracing has been used to estimate the volume of complex parts. This is accomplished by sending a set of parallel rays at the part; the fraction of rays that hit the part gives the cross-sectional area, and the average length that those rays are inside the part gives the average depth. Ray tracing has also often been used to determine proximity (including collision) for complex moving objects. This is usually done by sending "feeler" rays from the surfaces of objects to "see" what is nearby. Rays are also commonly used for mouse-based object selection to determine what object is seen in a pixel, and for projectile-object collision in games. OptiX can be used for any of those applications.

The common feature in ray tracing algorithms is that they compute the intersection points of 3D rays (an origin and a propagation direction) and a collection of 3D surfaces (the "model" or "scene"). In rendering applications, the optical properties of the point where the ray intersects the model determine what happens to the ray (for example, it might be reflected, absorbed or refracted). Other applications might not care about information other than where the intersection happens, or even if an intersection occurs at all. This variety of needs means it is desirable for OptiX to support a variety of ray-scene queries and user-defined behavior when rays intersect the scene.

One of ray tracing's nice features is that it is easy to support any geometric object that can be intersected with a 3D line. For example, it is straightforward to support spheres natively with no tessellation. Another nice feature is that ray tracing's execution is normally "sub-linear" in the number of objects—doubling the number of objects in the scene should less than double

the running time. This is accomplished by organizing the objects into an acceleration structure that can quickly reject whole groups of primitives as not candidates for intersection with any given ray. For static parts of the scene, this structure can be reused for the life of the application. For dynamic parts of the scene, OptiX supports rebuilding the acceleration structure when needed. The structure only queries the bounding box of any geometric objects it contains, so new types of primitives can be added and the acceleration structures will continue to work without modification, so long as the new primitives can provide a bounding box.

For graphics applications, ray tracing has advantages over rasterization. One of these is that general camera models are easy to support; the user can associate points on the screen with any direction they want, and there is no requirement that rays originate at the same point. Another advantage is that important optical effects such as reflection and refraction can be supported with only a few lines of code. Hard shadows are easy to produce with none of the artifacts typically associated with shadow maps, and soft shadows are not much harder. Furthermore, ray tracing can be added to more traditional graphics programs as a pass that produces a texture, letting the developer leverage the best of both worlds. For example, just the specular reflections could be computed by using points in the depth buffer as ray origins. There are a number of such "hybrid algorithms" that use both z-buffer and ray tracing techniques.

1.4   The RTX platform

The NVIDIA RTX platform combines ray tracing, deep learning and rasterization to enhance the creative process for content creators and developers, leveraging the RT cores on NVIDIA Turing GPUs. Software developers can use the OptiX API to take advantage of this platform. The OptiX AI denoiser uses NVIDIA Tensor Cores to accelerate ray tracing by reducing the time required to produce a noise-free image.

> **Note:** RTX mode is not supported by Kepler GPUs and requires Maxwell GPUs or later.

2   Programming model

The OptiX programming model consists of two halves: the host code and the GPU device programs. This chapter introduces the objects, programs, and variables that are defined in host code and used on the device.

2.1   Object model

OptiX is an object-based C API that implements a simple retained mode object hierarchy. This object-oriented host interface is augmented with programs that execute on the GPU. The main objects in the system are:

`Context`
   An instance of a running OptiX engine

`Program`
   A CUDA C function, compiled to NVIDIA PTX virtual assembly language

`Variables`
   A name used to pass data from C to OptiX programs

`Buffer`
   A multidimensional array that can be bound to a variable

`TextureSampler`
   One or more buffers bound with an interpolation mechanism

`Geometry`
   One or more primitives that a ray can be intersected with, such as triangles or other user-defined types

`Material`
   A set of programs executed when a ray intersects with the closest primitive or potentially closest primitive

`GeometryInstance`
   A binding between `Geometry` and `Material` objects.

`GroupNode`
   A set of objects arranged in a hierarchy

`GeometryGroup`
   A set of `GeometryInstance` objects

`TransformNode`
   A hierarchy node that geometrically transforms rays, so as to transform the geometric objects

`SelectorNode`
   A programmable hierarchy node that selects which children to traverse

`AccelerationStructure`
   An acceleration structure object that can be bound to a hierarchy node

---

`RemoteDeviceObj`
> A network connection for optional remote rendering.

These objects are created, destroyed, modified and bound with the C API and are further detailed in Chapter 3. The behavior of OptiX can be controlled by assembling these objects into any number of different configurations.

2.2   Component programs

The ray tracing pipeline provided by OptiX contains several programmable components. These programs are invoked on the GPU at specific points during the execution of a generic ray tracing algorithm. There are eight types of programs:

Ray generation programs (page 47)
> The entry point into the ray tracing pipeline, invoked by the system in parallel for each pixel, sample, or other user-defined work assignment

Exception programs (page 50)
> Exception handler, invoked for conditions such as stack overflow and other errors

Closest hit programs (page 52)
> Called when a traced ray finds the closest intersection point, such as for material shading

Any hit programs (page 53)
> Called when a traced ray finds a new, potentially closest, intersection point, such as for shadow computation

Intersection programs (page 56)
> Implements a ray-primitive intersection test, invoked during traversal

Bounding box programs (page 56)
> Computes a primitive's world space bounding box, called when the system builds a new acceleration structure over the geometry

Miss programs (page 55)
> Called when a traced ray misses all scene geometry

Visit programs (page 59)
> Called during traversal of a `Selector` node to determine the children a ray will traverse

The input language for these programs is PTX. The OptiX SDK also provides a set of wrapper classes and headers for use with the NVIDIA C Compiler (nvcc) that enable the use of CUDA C as a way of generating appropriate PTX.

These programs are further detailed in the "Programs" (page 41) chapter.

2.3   Variables

OptiX features a flexible and powerful variable system for communicating data to programs. When an OptiX program references a variable, there is a well-defined set of scopes that will be queried for a definition of that variable. This enables dynamic overrides of variable definitions based on which scopes are queried for definitions.

For example, a closest hit program may reference a variable called `color`. This program may then be attached to multiple `Material` objects, which are, in turn, attached to `GeometryInstance` objects. Variables in closest hit programs first look for definitions directly

attached to their `Program` object, followed by `GeometryInstance`, `Material` and `Context` objects, in that order. This enables a default color definition to exist on the `Material` object but specific instances using that material to override the default color definition.

See "Graph nodes" (page 22) for more information.

2.4 Execution model

Once all of these objects, programs and variables are assembled into a valid context, ray generation programs may be launched. Launches take dimensionality and size parameters and invoke the ray generation program a number of times equal to the specified size.

Once the ray generation program is invoked, a special semantic variable may be queried to provide a run-time index identifying the ray generation program invocation. For example, a common use case is to launch a two-dimensional invocation with a width and height equal to the size, in pixels, of an image to be rendered.

See "Launching a ray generation program" (page 48) for more information on launching ray generation programs from a context.

3 Host API

3.1 Context

An OptiX context provides an interface for controlling the setup and subsequent launch of the ray tracing engine. Contexts are created with the `rtContextCreate` function. A context object encapsulates all OptiX resources—textures, geometry, user-defined programs, etc. The destruction of a context, via the `rtContextDestroy` function, will clean up all of these resources and invalidate any existing handles to them.

The functions `rtContextLaunch1D`, `rtContextLaunch2D` and `rtContextLaunch3D` (collectively known as `rtContextLaunch`) serve as entry points to ray engine computation. The launch function takes an entry point parameter, discussed in , as well as one, two or three grid dimension parameters. The dimensions establish a logical computation grid. Upon a call to `rtContextLaunch`, any necessary preprocessing is performed and then the ray generation program associated with the provided entry point index is invoked once per computational grid cell. The launch precomputation includes state validation and, if necessary, acceleration structure generation and kernel compilation. Output from the launch is passed back via OptiX buffers, typically but not necessarily of the same dimensionality as the computation grid.

Listing 3.1

```
RTcontext context;
rtContextCreate( &context );
unsigned int entry_point = ...;
unsigned int width = ...;
unsigned int height = ...;

...    Set up context state and scene description

rtContextLaunch2D( context, entry_point, width, height );
rtContextDestroy( context );
```

While multiple contexts can be active at one time in limited cases, this is usually unnecessary as a single context object can leverage multiple hardware devices.

As of OptiX 4.0, mixed multi-GPU setups are available on all supported GPU architectures which are Kepler, Maxwell, Pascal, and Volta GPUs.

By default all compatible GPU devices in a system will be selected in an OptiX context when not explicitly using the function `rtContextSetDevices` to specify which devices should be made available. If incompatible devices are selected an error is returned from `rtContextSetDevices`.

In mixed GPU configurations, the kernel will be compiled for each streaming multiprocessor (SM) architecture, extending the initial start-up time.

For best performance, use multi-GPU configurations consisting of the same GPU type. Also prefer PCI-E slots in the system with the highest number of electrical PCI-E lanes (x16 Gen3 recommended).

On system configurations without NVLINK support, the board with the smallest VRAM amount will be the limit for on-device resources in the OptiX context. In homogeneous multi-GPU systems with NVLINK bridges and the driver running in the Tesla Compute Cluster (TCC) mode, OptiX will automatically use peer-to-peer access across the NVLINK connections to use the combined VRAM of the individual boards together which allows bigger scene sizes.

3.1.1   Entry points

Each context may have multiple computation entry points. A context entry point is associated with a single ray generation program as well as an exception program. The total number of entry points for a given context can be set with `rtContextSetEntryPointCount`. Each entry point's associated programs are set by `rtContextSetRayGenerationProgram` and `rtContextSetExceptionProgram` and are queried by `rtContextGetRayGenerationProgram` and `rtContextGetExceptionProgram`. Each entry point must be assigned a ray generation program before use; however, the exception program is an optional program that allows users to specify behavior upon various error conditions. The multiple entry point mechanism allows switching between multiple rendering algorithms as well as efficient implementation of techniques such as multi-pass rendering on a single OptiX context.

Listing 3.2

```
RTcontext context = ...;
rtContextSetEntryPointCount( context, 2 );

RTprogram pinhole_camera = ...;
RTprogram thin_lens_camera = ...;
RTprogram exception = ...;

rtContextSetRayGenerationProgram( context, 0, pinhole_camera );
rtContextSetRayGenerationProgram( context, 1, thin_lens_camera );

rtContextSetExceptionProgram( context, 0, exception );
rtContextSetExceptionProgram( context, 1, exception );
```

3.1.2   Ray types

OptiX supports the notion of ray types, which is useful to distinguish between rays that are traced for different purposes. For example, a renderer might distinguish between rays used to compute color values and rays used exclusively for determining visibility of light sources (*shadow rays*). Proper separation of such conceptually different ray types not only increases program modularity, but also enables OptiX to operate more efficiently.

Both the number of different ray types as well as their behavior is entirely defined by the application. The number of ray types to be used is set with `rtContextSetRayTypeCount`.

The following properties may differ among ray types:

- The ray payload

- The closest-hit (page 52) program of each individual material

- The any-hit (page 53) program of each individual material

- The miss (page 55) program

The ray payload is an arbitrary user-defined data structure associated with each ray. This is commonly used, for example, to store a result color, the ray's recursion depth, a shadow attenuation factor, and so on. It can be regarded as the result a ray delivers after having been traced, but it can also be used to store and propagate data between ray generations during recursive ray tracing.

The closest-hit (page 52) and any-hit (page 53) programs assigned to materials correspond roughly to shaders in conventional rendering systems: they are invoked when an intersection between a ray and a geometric primitive is found. Since those programs are assigned to materials per ray type, not all ray types must define behavior for both program types. See Closest-hit programs (page 52) and Any-hit programs (page 53) for a more detailed discussion of material programs.

The miss (page 55) program is executed when a traced ray is determined to not hit any geometry. A miss program could, for example, return a constant sky color or sample from an environment map.

As an example of how to make use of ray types, a Whitted-style recursive ray tracer might define the ray types listed in Table 1:

Use	Radiance	Shadow
Payload	RadiancePL	ShadowPL
Closest hit	Compute color, keep track of recursion depth	—
Any hit	—	Compute shadow attenuation and terminate ray if opaque
Miss	Environment map lookup	—

Table 1 – Example ray types

The ray payload data structures in the above example might look as follows:

Listing 3.3

```
struct RadiancePL {
    float3 color;                   Payload for ray type 0: radiance rays
    int    recursion_depth;
};

struct ShadowPL {
    float attenuation;    Payload for ray type 1: shadow rays
};
```

Upon a call to `rtContextLaunch`, the ray generation program traces radiance rays into the scene, and writes the delivered results (found in the `color` field of the payload) into an output buffer for display:

Listing 3.4

```
RadiancePL payload;
payload.color = make_float3( 0.f, 0.f, 0.f );
payload.recursion_depth = 0;    Initialize recursion depth

Ray ray = ...    Some camera code creates the ray

ray.ray_type = 0;    Make this a radiance ray

rtTrace( top_object, ray, payload );

writeOutput( payload.color );    Write result to output buffer
```

A primitive intersected by a radiance ray would execute a closest-hit program which computes the ray's color and potentially traces shadow rays and reflection rays. The shadow ray part is shown in the following code snippet:

Listing 3.5

```
ShadowPL shadow_payload;
shadow_payload.attenuation = 1.0f;    Initialize to visible

Ray shadow_ray = ...    Create a ray to light source

shadow_ray.ray_type = 1;    Make this a shadow ray
rtTrace( top_object, shadow_ray, shadow_payload );

float3 rad =                                      Attenuate incoming light ("light"
    light.radiance * shadow_payload.attenuation;  is some user-defined variable
                                                  describing the light source)
```

| `payload.color += rad;` | Add the contribution to the current radiance ray's payload (assumed to be declared as "payload") |

To properly attenuate shadow rays, all materials use an [any-hit](page 53) program which adjusts the attenuation and terminates ray traversal. The following code sets the attenuation to zero, assuming an opaque material:

Listing 3.6

| `shadow_payload.attenuation = 0;` | Assume opaque material |

| `rtTerminateRay();` | It won't get any darker, so terminate |

3.1.3   Global state

Aside from ray type and entry point counts, there are several other global settings encapsulated within OptiX contexts.

Each context holds a number of attributes that can be queried and set using `rtContextGetAttribute` and `rtContextSetAttribute`. For example, the amount of memory an OptiX context has allocated on the host can be queried by specifying `RT_CONTEXT_ATTRIBUTE_USED_HOST_MEMORY` as attribute parameter.

To support recursion, OptiX uses a small stack of memory associated with each thread of execution. `rtContextGetStackSize` and `rtContextSetStackSize` enable setting and querying the size of this stack. The stack size should be set with care as unnecessarily large stacks will result in performance degradation while overly small stacks will cause overflows within the ray engine. Stack overflow errors can be handled with user defined *exception programs*.

The `rtContextSetPrint` functions are used to enable C-style `printf` printing from within OptiX programs, allowing these programs to be more easily debugged. The CUDA C function `rtContextSetPrintEnabled` turns on or off printing globally while `rtContextSetPrintLaunchIndex` toggles printing for individual computation grid cells. Print statements have no adverse effect on performance while printing is globally disabled, which is the default behavior.

Print requests are buffered in an internal buffer, the size of which can be specified with `rtContextSetPrintBufferSize`. Overflow of this buffer will cause truncation of the output stream. The output stream is printed to the standard output after all computation has completed but before `rtContextLaunch` has returned.

Listing 3.7

```
RTcontext context = ...;
rtContextSetPrintEnabled( context, 1 );
rtContextSetPrintBufferSize( context, 4096 );
```

Within an OptiX program, the `rtPrintf` function works similarly to C's `printf`. Each invocation of `rtPrintf` will be atomically deposited into the print output buffer, but separate invocations by the same thread or by different threads will be interleaved arbitrarily.

Listing 3.8

```
rtDeclareVariable(uint2, launch_idx, rtLaunchIndex, );

RT_PROGRAM void any_hit()
{
    rtPrintf( "Hello from index %u, %u!\n", launch_idx.x, launch_idx.y );
}
```

The context also serves as the outermost scope for OptiX variables. Variables declared via `rtContextDeclareVariable` are available to all OptiX objects associated with the given context. To avoid name conflicts, existing variables may be queried with either `rtContextQueryVariable` (by name) or `rtContextGetVariable` (by index), and removed with `rtContextRemoveVariable`.

`rtContextValidate` can be used at any point in the setup process to check the state validity of a context and all of its associated OptiX objects. This will include checks for the presence of necessary programs (for example, an intersection program for a geometry node), invalid internal state such as unspecified children in graph nodes and the presence of variables referred to by all specified programs. Validation is always implicitly performed upon a context launch.

`rtContextSetTimeoutCallback` specifies a callback function of type `RTtimeoutcallback` that is called at a specified maximum frequency from OptiX API calls that can run long, such as acceleration structure builds, compilation, and kernel launches. This allows the application to update its interface or perform other tasks. The callback function may also ask OptiX to cease its current work and return control to the application. This request is complied with as soon as possible. Output buffers expected to be written to by an `rtContextLaunch` are left in an undefined state, but otherwise OptiX tracks what tasks still need to be performed and resumes cleanly in subsequent API calls.

Listing 3.9

```
int timeout_callback()
{
    update_gui();
    return check_gui_status();
}
...
```
An `RTtimeoutcallback`: Return 1 to ask for abort, 0 to continue.

```
rtContextSetTimeoutCallback(
    context, timeout_
callback, 0.1 );
```
Call timeout_callback() at most once every 100 ms.

`rtContextGetErrorString` can be used to get a description of any failures occurring during context state setup, validation, or launch execution.

3.1.4   Enabling RTX mode

As of OptiX version 5.2, RTX mode can be enabled to take advantage of RT Cores, accelerating ray tracing by computing traversal and triangle intersection in hardware.

RTX mode is not enabled by default. RTX mode can be enabled with the
`RT_GLOBAL_ATTRIBUTE_ENABLE_RTX` attribute using `rtGlobalSetAttribute` when creating the
OptiX context. However, certain features of OptiX will not be available.

3.2 Buffers

OptiX uses buffers to pass data between the host and the device. Buffers are created by the
host prior to invocation of `rtContextLaunch` using the `rtBufferCreate` function. This
function also sets the buffer type as well as optional flags. The type and flags are specified as a
bitwise OR combination.

The buffer type determines the direction of data flow between host and device. Its options are
enumerated by `RTbuffertype`:

`RT_BUFFER_INPUT`

> Only the host may write to the buffer. Data is transferred from host to device and device
> access is restricted to be read-only.

`RT_BUFFER_OUTPUT`

> The converse of `RT_BUFFER_INPUT`. Only the device may write to the buffer. Data is
> transferred from device to host.

`RT_BUFFER_INPUT_OUTPUT`

> Allows read-write access from both the host and the device.

`RT_BUFFER_PROGRESSIVE_STREAM`

> The automatically updated output of a progressive launch. Can be streamed efficiently
> over network connections. (See "Progressive launches" (page 37).)

Buffer flags specify certain buffer characteristics and are enumerated by the fields of
`RTbufferflag`:

`RT_BUFFER_GPU_LOCAL`

> Can only be used in combination with `RT_BUFFER_INPUT_OUTPUT`. This restricts the host
> to write operations as the buffer is not copied back from the device to the host. The device
> is allowed read-write access. However, writes from multiple devices are not coherent, as a
> separate copy of the buffer resides on each device.

`RT_BUFFER_LAYERED`

> If `RT_BUFFER_LAYERED` flag is set, buffer depth specifies the number of layers, not the
> depth of a 3D buffer, when it is used as a texture buffer.

`RT_BUFFER_CUBEMAP`

> If `RT_BUFFER_CUBEMAP` flag is set, buffer depth specifies the number of cube faces, not the
> depth of a 3D buffer.

Before using a buffer, its size, dimensionality and element format must be specified. The
format can be set and queried with `rtBufferSetFormat` and `rtBufferGetFormat`. Format
options are enumerated by the `RTformat` type. Formats exist for C and CUDA C data types
such as `unsigned int` and `float3`. Buffers of arbitrary elements can be created by choosing
the format `RT_FORMAT_USER` and specifying an element size with the
`rtBufferSetElementSize` function. The size of the buffer is set with `rtBufferSetSize1D`,
`rtBufferSetSize2D`, and `rtBufferSetSize3D`, which also specify the dimensionality
implicitly. Functions `rtBufferGetMipLevelSize1D`, `rtBufferGetMipLevelSize2D`, and

`rtBufferGetMipLevelSize3D` can be used to get the size of a level image in the mipmap, given its level number.

Listing 3.10

```
RTcontext context = ...;
RTbuffer buffer;
typedef struct { float r; float g; float b; } rgb;

rtBufferCreate( context, RT_BUFFER_INPUT_OUTPUT, &buffer );
rtBufferSetFormat( RT_FORMAT_USER );
rtBufferSetElementSize( sizeof(rgb) );
rtBufferSetSize2D( buffer, 512, 512 );
```

Host access to the data stored within a buffer is performed with the `rtBufferMap` function. This function returns a pointer to a one dimensional array representation of the buffer data. All buffers must be unmapped via `rtBufferUnmap` before context validation will succeed.

Listing 3.11

```
unsigned int width, height;
rtBufferGetSize2D( buffer, &width, &height );        Using the buffer created above

void* data;
rtBufferMap( buffer, &data );

rgb* rgb_data = (rgb*)data;
for( unsigned int i = 0; i < width*height; ++I ) {
    rgb_data[i].r = rgb_data[i].g = rgb_data[i].b =0.0f;
}
rtBufferUnmap( buffer );
```

`rtBufferMapEx` and `rtBufferUnmapEx` set the contents of a mip mapped texture buffer.

Listing 3.12

```
unsigned int width, height;
rtBufferGetMipLevelSize2D(                           Using the buffer created above
   buffer, &width, &height, level+1 );

rgb *dL, *dNextL;
rtBufferMapEx( buffer, RT_BUFFER_MAP_READ_WRITE, level, 0, &dL );
rtBufferMapEx( buffer, RT_BUFFER_MAP_READ_WRITE, level+1, 0, &dNextL );

unsigned int width2 = width*2;

for ( unsigned int y = 0; y < height; ++y ) {
   for ( unsigned int x = 0; x < width; ++x ) {
      dNextL[x+width*y] = 0.25f *
         (dL[x*2+width2*y*2] +
```

```
        dL[x*2+1+width2*y*2] +
        dL[x*2+width2*(y*2+1)] +
        dL[x*2+1+width2*(y*2+1)]);
    }
}
rtBufferUnmapEx( buffer, level );
rtBufferUnmapEx( buffer, level+1 );
```

Access to buffers within OptiX programs uses a simple array syntax. The two template arguments in the declaration below are the element type and the dimensionality, respectively.

Listing 3.13

```
rtBuffer<rgb, 2> buffer;
...
uint2 index = ...;
float r = buffer[index].r;
```

3.2.1   Buffers of buffer ids

Beginning in OptiX 3.5, buffers may contain IDs to buffers. From the host side, an input buffer is declared with format `RT_FORMAT_BUFFER_ID`. The buffer is then filled with buffer IDs obtained through the use of either `rtBufferGetId` or `BufferObj::getId`. A special sentinel value, `RT_BUFFER_ID_NULL`, can be used to distinguish between valid and invalid buffer IDs. `RT_BUFFER_ID_NULL` will never be returned as a valid buffer ID.

The following example that creates two input buffers; the first contains the data, and the second contains the buffer IDs.

Listing 3.14

```
Buffer inputBuffer0 =
    context->createBuffer( RT_BUFFER_INPUT, RT_FORMAT_INT, 3 );
Buffer inputBuffers =
    context->createBuffer( RT_BUFFER_INPUT, RT_FORMAT_BUFFER_ID, 1);
int* buffers = static_cast<int*>(inputBuffers->map());
buffers[0] = inputBuffer0->getId();
inputBuffers->unmap();
```

From the device side, buffers of buffer IDs are declared using `rtBuffer` with a template argument type of `rtBufferId`. The identifiers stored in the buffer are implicitly cast to buffer handles when used on the device. This example creates a one dimensional buffer whose elements are themselves one dimensional buffers that contain integers.

Listing 3.15

```
rtBuffer<rtBufferId<int,1>, 1> input_buffers;
```

Accessing the buffer is done the same way as with regular buffers:

Listing 3.16

`int value = input_buffers[buf_index][0];`	Grab the first element of the first buffer in "input_buffers"

The size of the buffer can also be queried to loop over the contents:

Listing 3.17

```
for(size_t i = 0; k < input_buffers.size(); ++i)
    result += input_buffers[i];
```

Buffers may nest arbitrarily deeply, though there is memory access overhead per nesting level. Multiple buffer lookups may be avoided by using references or copies of the `rtBufferId`.

Listing 3.18

```
rtBuffer<rtBufferId<rtBufferId<int,1>, 1>, 1> input_buffers3;
...
rtBufferId<int,1>& buffer = input_buffers[buf_index1][buf_index2];
size_t size = buffer.size();

for(size_t i = 0; i < size; ++i)
    value += buffer[i];
```

Currently only non-interop buffers of type `RT_BUFFER_INPUT` may contain buffer IDs and they may only contain IDs of buffers that match in element format and dimensionality, though they may have varying sizes.

The `RTbuffer` object associated with a given buffer ID can be queried with the function `rtContextGetBufferFromId` or if using the C++ interface, `ContextObj::getBufferFromId`.

In addition to storing buffer IDs in other buffers, you can store a buffer ID in an arbitrary struct, in an `RTvariable`, or as a data member in the ray payload. A buffer ID can also be passed as an argument to a callable program. An `rtBufferId` object can be constructed using the buffer ID as a constructor argument.

Listing 3.19

```
rtDeclareVariable(int, id, ,);
rtDeclareVariable(int, index, ,);
...
int value = rtBufferId<int,1>(id)[index];
```

An example of passing to a callable program:

Listing 3.20

```
#include <OptiX_world.h>
using namespace OptiX;

struct BufInfo {
```

```
    int index;
    rtBufferId<int, 1> data;
};


rtCallableProgram(int, getValue, (BufInfo));


RT_CALLABLE_PROGRAM
int getVal( BufInfo bufInfo )
{
    return bufInfo.data[bufInfo.index];
}


rtBuffer<int,1> result;
rtDeclareVariable(BufInfo, buf_info, ,);

RT_PROGRAM void bindlessCall()
{
    int value = getValue(buf_info);
    result[0] = value;
}
```

Note that because `rtCallableProgram` and `rtDeclareVariable` are macros, typedefs or structs should be used instead of using the templated type directly in order to work around the C preprocessor's limitations.

Listing 3.21

```
typedef rtBufferId<int,1> RTB;
rtDeclareVariable(RTB, buf, ,);
```

There is a definition for `rtBufferId` in `OptiXpp_namespace.h` that mirrors the device side declaration to enable declaring types that can be used in both host and device code.

Here is an example of the use of the BufInfo struct from the host side:

Listing 3.22

```
BufInfo buf_info;
buf_info.index = 0;
buf_info.data = rtBufferId<int,1>(inputBuf0->getId());
context["buf_info"]->setUserData(sizeof(buf_info), &buf_info);
```

3.3   Textures

OptiX textures provide support for common texture mapping functionality including texture filtering, various wrap modes, and texture sampling. `rtTextureSamplerCreate` is used to create texture objects. Each texture object is associated with one or more buffers containing the texture data. The buffers may be 1D, 2D or 3D and can be set with `rtTextureSamplerSetBuffer`.

`rtTextureSamplerSetFilteringModes` sets the filtering methods for minification, magnification and mipmapping. Wrapping for texture coordinates outside of the range [0.0,1.0] is specified per-dimension with `rtTextureSamplerSetWrapMode`.

The maximum anisotropy for a given texture is set with `rtTextureSamplerSetMaxAnisotropy`. This value will be clamped to the range [1.0,16.0].

`rtTextureSamplerSetReadMode` specifies that texture values are converted to normalized float values with a `readmode` parameter of `RT_TEXTURE_READ_NORMALIZED_FLOAT`.

Listing 3.23

```
RTcontext context = ...;
RTbuffer tex_buffer = ...;    2D buffer
RTtexturesampler tex_sampler;
rtTextureSamplerCreate( context, &tex_sampler );
rtTextureSamplerSetWrapMode( tex_sampler, 0, RT_WRAP_CLAMP_TO_EDGE);
rtTextureSamplerSetWrapMode( tex_sampler, 1, RT_WRAP_CLAMP_TO_EDGE);
rtTextureSamplerSetFilteringModes(
    tex_sampler, RT_FILTER_LINEAR, RT_FILTER_LINEAR, RT_FILTER_NONE );
rtTextureSamplerSetIndexingMode(
    tex_sampler, RT_TEXTURE_INDEX_NORMALIZED_COORDINATES );
rtTextureSamplerSetReadMode(
    tex_sampler, RT_TEXTURE_READ_NORMALIZED_FLOAT );
rtTextureSamplerSetMaxAnisotropy( tex_sampler, 1.0f );
rtTextureSamplerSetBuffer( tex_sampler, 0, 0, tex_buffer );
```

As of version 3.9, OptiX supports cube, layered, and mipmapped textures using new API calls `rtBufferMapEx`, `rtBufferUnmapEx`, `rtBufferSetMipLevelCount`.[1] Layered textures are equivalent to CUDA layered textures and OpenGL texture arrays. They are created by calling `rtBufferCreate` with `RT_BUFFER_LAYERED` and cube maps by passing `RT_BUFFER_CUBEMAP`. In both cases the buffer's depth dimension is used to specify the number of layers or cube faces, not the depth of a 3D buffer.

OptiX programs can access texture data with CUDA C's built-in `tex1D`, `tex2D` and `tex3D` functions.

Listing 3.24

```
rtTextureSampler<uchar4, 2, cudaReadModeNormalizedFloat> t;
...
float2 tex_coord = ...;
float4 value = tex2D( t, tex_coord.x, tex_coord.y );
```

As of version 3.0, OptiX supports *bindless textures*. Bindless textures allow OptiX programs to reference textures without having to bind them to specific variables. This is accomplished through the use of *texture IDs*.

Using bindless textures, it is possible to dynamically switch between multiple textures without the need to explicitly declare all possible textures in a program and without having to manually implement switching code. The set of textures being switched on can have

---

[1]`rtTextureSamplerSetArraySize` and `rtTextureSamplerSetMipLevelCount` are deprecated.

varying attributes, such as wrap mode, and varying sizes, providing increased flexibility over texture arrays.

To obtain a device handle from an existing texture sampler, `rtTextureSamplerGetId` can be used:

Listing 3.25

```
RTtexturesampler tex_sampler = ...;
int tex_id;
rtTextureSamplerGetId( tex_sampler, &tex_id );
```

A texture ID value is immutable and is valid until the destruction of its associated texture sampler. Make texture IDs available to OptiX programs by using input buffers or OptiX variables:

Listing 3.26

```
RTbuffer tex_id_buffer = ...;    1D buffer

unsigned int index = ...;

void* tex_id_data;
rtBufferMap( tex_id_buffer, &tex_id_data );
((int*)tex_id_data)[index] = tex_id;
rtBufferUnmap( tex_id_buffer );
```

Similar to CUDA C's texture functions, OptiX programs can access textures in a bindless way with `rtTex1D rtTex2D`, and `rtTex3D` functions:

Listing 3.27

```
rtBuffer<int, 1> tex_id_buffer;
unsigned int index = ...;
int tex_id = tex_id_buffer[index];
float2 tex_coord = ...;
float4 value = rtTex2D<float4>( tex_id, tex_coord.x, tex_coord.y );
```

Textures may also be sampled by providing a level of detail for mip mapping or gradients for anisotropic filtering. An integer layer number is required for layered textures (arrays of textures):

Listing 3.28

```
float4 v;
if ( mip_mode == MIP_DISABLE )
   v = rtTex2DLayeredLod<float4>( tex, uv.x, uv.y, tex_layer );
else if ( mip_mode == MIP_LEVEL )
   v = rtTex2DLayeredLod<float4>( tex, uv.x, uv.y, tex_layer, lod );
else if ( mip_mode == MIP_GRAD )
```

```
    v = rtTex2DLayeredGrad<float4>(
            tex, uv.x, uv.y, tex_layer, dpdx, dpdy );
```

3.4 Graph nodes

When a ray is traced from a program using the `rtTrace` function, a node is given that specifies the root of the graph. The host application creates this graph by assembling various types of nodes provided by the OptiX API. The basic structure of the graph is a hierarchy, with nodes describing geometric objects at the bottom, and collections of objects at the top.

The graph structure is not meant to be a scene graph in the classical sense. Instead, it serves as a way of binding different programs or actions to portions of the scene. Since each invocation of `rtTrace` specifies a root node, different trees or subtrees may be used. For example, shadowing objects or reflective objects may use a different representation—for performance or for artistic effect.

Graph nodes are created via `rt*Create` calls, which take the Context as a parameter. Since these graph node objects are owned by the context, rather than by their parent node in the graph, a call to `rt*Destroy` will delete that object's variables, but not do any reference counting or automatic freeing of its child nodes.

Figure 3.1 shows an example of what a graph might look like. The following sections will describe the individual node types.
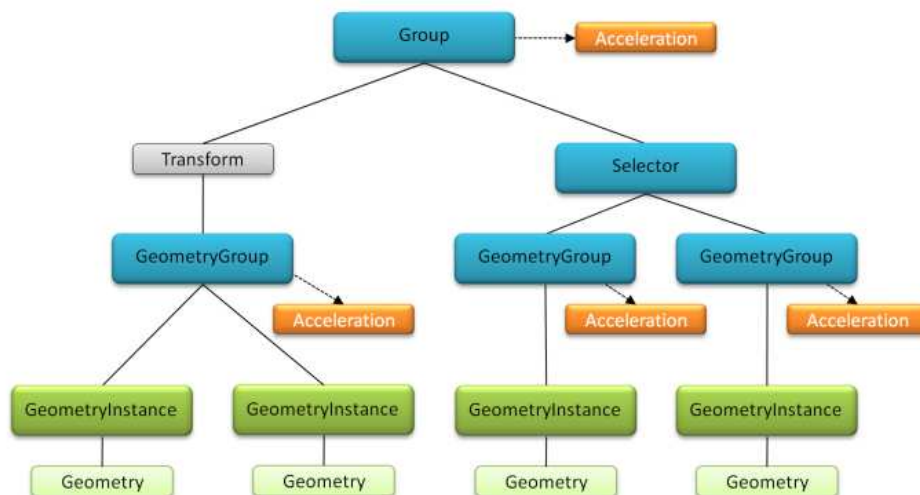
Fig. 3.1 – A sample graph

Table 2 indicates which nodes can be children of other nodes including association with acceleration structure and material nodes.

Parent node type	Child node types	Associated node types
Geometry Acceleration	*none*	Material
GeometryInstance	Geometry	Material
GeometryGroup	GeometryInstance	Acceleration
Transform Selector	GeometryGroup Transform	*node*
Group	GeometryGroup	Acceleration

Table 2 – Parent nodes and the types of nodes allowed as children

3.4.1 Geometry

A geometry node is the fundamental node to describe a geometric object: a collection of user-defined primitives against which rays can be intersected. The number of primitives contained in a geometry node is specified using `rtGeometrySetPrimitiveCount`.

To define the primitives, an intersection program is assigned to the geometry node using `rtGeometrySetIntersectionProgram`. The input parameters to an intersection program are a primitive index and a ray, and it is the program's job to return the intersection between the two. In combination with program variables, this provides the necessary mechanisms to define any primitive type that can be intersected against a ray. A common example is a triangle mesh, where the intersection program reads a triangle's vertex data out of a buffer (passed to the program via a variable) and performs a ray-triangle intersection.

In order to build an acceleration structure over arbitrary geometry, it is necessary for OptiX to query the bounds of individual primitives. For this reason, a separate bounds program must be provided using `rtGeometrySetBoundingBoxProgram`. This program simply computes bounding boxes of the requested primitives, which are then used by OptiX as the basis for acceleration structure construction.

The following example shows how to construct a geometry object describing a sphere, using a single primitive. The intersection and bounding box program are assumed to depend on a single parameter variable specifying the sphere radius:

Listing 3.29

```
RTgeometry geometry;
RTvariable variable;


rtGeometryCreate( context, &geometry );
rtGeometrySetPrimitiveCount( geometry, 1 );
rtGeometrySetIntersectionProgram(
    geometry, sphere_intersection );
rtGeometrySetBoundingBoxProgram(
    geometry, sphere_bounds );
```

Set up geometry object.

```
rtGeometryDeclareVariable(
    geometry, "radius", &variable );    Declare and set the radius variable.
rtVariableSet1f( variable, 10.0f );
```

3.4.2 Material

A material encapsulates the actions that are taken when a ray intersects a primitive associated with a given material. Examples for such actions include: computing a reflectance color, tracing additional rays, ignoring an intersection, and terminating a ray. Arbitrary parameters can be provided to materials by declaring program variables.

Two types of programs may be assigned to a material, closest-hit programs and any-hit programs. The two types differ in when and how often they are executed. The closest-hit program, which is similar to a shader in a classical rendering system, is executed at most once per ray, for the closest intersection of a ray with the scene. It typically performs actions that involve texture lookups, reflectance color computations, light source sampling, recursive ray tracing, and so on, and stores the results in a ray payload data structure.

The any-hit program is executed for each potential closest intersection found during ray traversal. The intersections for which the program is executed may not be ordered along the ray, but eventually all intersections of a ray with the scene can be enumerated if required (by calling `rtIgnoreIntersection` on each of them). Typical uses of the any-hit program include early termination of shadow rays (using `rtTerminateRay`) and binary transparency, for example, by ignoring intersections based on a texture lookup.

It is important to note that both types of programs are assigned to materials *per ray type*, which means that each material can actually hold more than one closest-hit or any-hit program. This is useful if an application can identify that a certain kind of ray only performs specific actions. For example, a separate ray type may be used for shadow rays, which are only used to determine binary visibility between two points in the scene. In this case, a simple any-hit program attached to all materials under that ray type index can immediately terminate such rays, and the closest-hit program can be omitted entirely. This concept allows for highly efficient specialization of individual ray types.

The closest-hit program is assigned to the material by calling `rtMaterialSetClosestHitProgram`, and the any-hit program is assigned with `rtMaterialSetAnyHitProgram`. If a program is omitted, an empty program is the default.

3.4.3 GeometryInstance

A geometry instance represents a coupling of a single geometry node with a set of materials. The geometry object the instance refers to is specified using `rtGeometryInstanceSetGeometry`. The number of materials associated with the instance is set by `rtGeometryInstanceSetMaterialCount`, and the individual materials are assigned with `rtGeometryInstanceSetMaterial`. The number of materials that must be assigned to a geometry instance is determined by the highest material index that may be reported by an intersection program of the referenced geometry.

Note that multiple geometry instances are allowed to refer to a single geometry object, enabling instancing of a geometric object with different materials. Likewise, materials can be reused between different geometry instances.

This example configures a geometry instance so that its first material index is `mat_phong` and the second one is `mat_diffuse`, both of which are assumed to be `RTmaterial` objects with appropriate programs assigned. The instance is made to refer to the `RTgeometry` object `triangle_mesh`.

Listing 3.30

```
RTgeometryinstance ginst;

rtGeometryInstanceCreate( context, &ginst );
rtGeometryInstanceSetGeometry( ginst, triangle_mesh );

rtGeometryInstanceSetMaterialCount( ginst, 2 );
rtGeometryInstanceSetMaterial( ginst, 0, mat_phong );
rtGeometryInstanceSetMaterial( ginst, 1, mat_diffuse);
```

3.4.4   GeometryGroup

A geometry group is a container for an arbitrary number of geometry instances. The number of contained geometry instances is set using `rtGeometryGroupSetChildCount`, and the instances are assigned with `rtGeometryGroupSetChild`. Each geometry group must also be assigned an acceleration structure using `rtGeometryGroupSetAcceleration`. (See "Acceleration structures" (page 27).)

The minimal sample use case for a geometry group is to assign it a single geometry instance:

Listing 3.31

```
RTgeometrygroup geomgroup;

rtGeometryGroupCreate( context, &geomgroup );
rtGeometryGroupSetChildCount( geomgroup, 1 );
rtGeometryGroupSetChild( geomgroup, 0, geometry_instance );
```

Multiple geometry groups are allowed to share children, that is, a geometry instance can be a child of more than one geometry group.

3.4.5 Group

A group represents a collection of higher level nodes in the graph. They are used to compile the graph structure which is eventually passed to `rtTrace` for intersection with a ray.

A group can contain an arbitrary number of child nodes, which must themselves be of type `RTgroup`, `RTgeometrygroup`, `RTtransform`, or `RTselector`. The number of children in a group is set by `rtGroupSetChildCount`, and the individual children are assigned using `rtGroupSetChild`. Every group must also be assigned an acceleration structure via `rtGroupSetAcceleration`.

A common use case for groups is to collect several geometry groups which dynamically move relative to each other. The individual position, rotation, and scaling parameters can be represented by `Transform` nodes, so the only acceleration structure that needs to be rebuilt between calls to `rtContextLaunch` is the one for the top level group. This will usually be much cheaper than updating acceleration structures for the entire scene.

Note that the children of a group can be shared with other groups, that is, each child node can also be the child of another group (or of any other graph node for which it is a valid child). This allows for very flexible and lightweight instancing scenarios, especially in combination with shared acceleration structures. (See "Acceleration structures" (page 27).)

3.4.6 Transform

A `Transform` node is used to represent a projective transformation of its underlying scene geometry. The transform must be assigned exactly one child of type `RTgroup`, `RTgeometrygroup`, `RTtransform`, or `RTselector`, using `rtTransformSetChild`. That is, the nodes below a transform may simply be geometry in the form of a geometry group, or a whole new subgraph of the scene.

The transformation itself is specified by passing a 4x4 floating point matrix (specified as a 16-element one-dimensional array) to `rtTransformSetMatrix`. Conceptually, it can be seen as if the matrix were applied to all the underlying geometry. However, the effect is instead achieved by transforming the rays themselves during traversal. This means that OptiX does not rebuild any acceleration structures when the transform changes.

This example shows how a `Transform` object with a simple translation matrix is created:

Listing 3.32

```
RTtransform transform;
const float x=10.0f, y=20.0f, z=30.0f;

const float m[16] = {
    1, 0, 0, x,
    0, 1, 0, y,              Matrices are row-major.
    0, 0, 1, z,
    0, 0, 0, 1
};

rtTransformCreate( context, &transform );
rtTransformSetMatrix( transform, 0, m, 0 );
```

Note that a `Transform` child node may be shared with other graph nodes. That is, a child node of a `Transform` may be a child of another node at the same time. This is often useful for instancing geometry.

`Transform` nodes should be used sparingly as they cost performance during ray tracing. In particular, it is highly recommended for node graphs to not exceed a single level of transform depth.

3.4.7 Selector

A `Selector` is similar to a group in that it is a collection of higher level graph nodes. The number of nodes in the collection is set by `rtSelectorSetChildCount`, and the individual children are assigned with `rtSelectorSetChild`. Valid child types are `RTgroup`, `RTgeometrygroup`, `RTtransform`, and `RTselector`.

The main difference between selectors and groups is that selectors do not have an acceleration structure associated with them. Instead, a visit (page 59) program is specified with `rtSelectorSetVisitProgram`. This program is executed every time a ray encounters the selector node during graph traversal. The program specifies which children the ray should continue traversal through by calling `rtIntersectChild`.

A typical use case for a selector is dynamic (per-ray) level of detail: an object in the scene may be represented by a number of geometry nodes, each containing a different level of detail version of the object. The geometry groups containing these different representations can be assigned as children of a selector. The visit program can select which child to intersect using any criterion (for example, based on the footprint or length of the current ray), and ignore the others.

As for groups and other graph nodes, child nodes of a selector can be shared with other graph nodes to allow flexible instancing.

> **Note:** Selector nodes are deprecated in RTX mode. Future updates to RTX mode will provide a mechanism to support most of the use cases that required Selector nodes. See Enabling RTX mode (page 14).

3.5   Acceleration structures

Acceleration structures are an important tool for speeding up the traversal and intersection queries for ray tracing, especially for large scene databases. Most successful acceleration structures represent a hierarchical decomposition of the scene geometry. This hierarchy is then used to quickly cull regions of space not intersected by the ray.

There are different types of acceleration structures, each with their own advantages and drawbacks. Furthermore, different scenes require different kinds of acceleration structures for optimal performance (for example, the difference between static and dynamic scenes, generic primitives and triangles, and so on). The most common trade-off is construction speed vs. ray tracing performance, but other factors such as memory consumption can play a role as well.

No single type of acceleration structure is optimal for all scenes. To allow an application to balance the trade-offs, OptiX lets you choose between several kinds of supported structures. You can even mix and match different types of acceleration structures within the same node graph.

3.5.1   Acceleration objects in the node graph

Acceleration structures are individual API objects in OptiX, called `RTacceleration`. Once an acceleration object is created with `rtAccelerationCreate`, it is assigned to either a group (using `rtGroupSetAcceleration`) or a geometry group (using `rtGeometryGroupSetAcceleration`). Every group and geometry group in the node graph needs to have an acceleration object assigned for ray traversal to intersect those nodes.

This example creates a geometry group and an acceleration structure and connects the two:

Listing 3.33

```
RTgeometrygroup geomgroup;
RTacceleration accel;

rtGeometryGroupCreate( context, &geomgroup );
rtAccelerationCreate( context, &accel );
rtGeometryGroupSetAcceleration( geomgroup, accel );
```

By making use of groups and geometry groups when assembling the node graph, the application has a high level of control over how acceleration structures are constructed over the scene geometry. If one considers the case of several geometry instances in a scene, there are a number of ways they can be placed in groups or geometry groups to fit the application's use case.

For example, Figure 3.2 (page 29) places all the geometry instances in a single geometry group. An acceleration structure on a geometry group will be constructed over the individual primitives defined by the collection of child geometry instances. This will allow OptiX to build an acceleration structure which is as efficient as if the geometries of the individual instances had been merged into a single object.
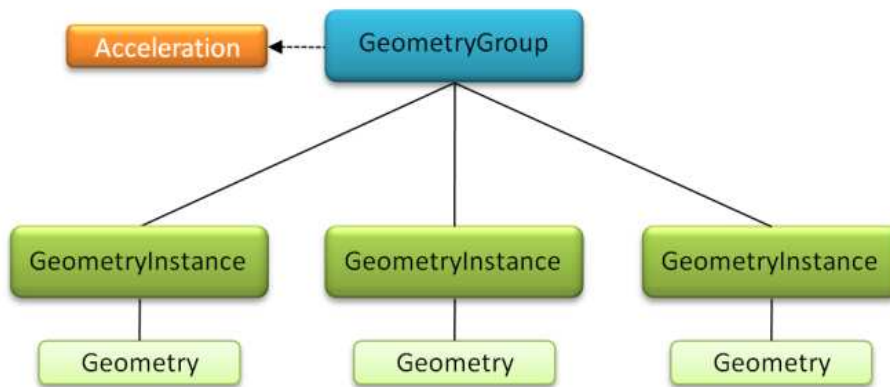
Fig. 3.2 – Multiple geometry instances in a geometry group

A different approach to managing multiple geometry instances is shown in Figure 3.3. Each instance is placed in its own geometry group, so that there is a separate acceleration structure for each instance. The resulting collection of geometry groups is aggregated in a top level group, which itself has an acceleration structure. Acceleration structures on groups are constructed over the bounding volumes of the child nodes. Because the number of child nodes is usually relatively low, high level structures are typically quick to update. The advantage of this approach is that when one of the geometry instances is modified, the acceleration structures of the other instances need not be rebuilt. However, because higher level acceleration structures introduce an additional level of complexity and are built only on the coarse bounds of their group's children, the graph in Figure 3.3 will likely not be as efficient to traverse as the one in Figure 3.2. Again, this is a trade-off the application needs to balance, in this case, by considering how frequently individual geometry instances will be modified.

Fig. 3.3 – Multiple geometry instances, each in a separate geometry group

3.5.2 Acceleration structure builders

An `RTacceleration` has a *builder*. The builder is responsible for collecting input geometry (in most cases, this geometry is the bounding boxes created by geometry nodes' bounding box programs) and computing a data structure that allows for accelerated ray-scene intersection

query. Builders are not application-defined programs. Instead, the application chooses an appropriate builder from Table 3.

Builder	Description
Trbvh	The Trbvh[2] . builder performs a very fast GPU-based BVH build. Its ray tracing performance is usually within a few percent of SBVH, yet its build time is generally the fastest. This builder should be strongly considered for all data sets. Trbvh uses a modest amount of extra memory beyond that required for the final BVH. When the extra memory is not available on the GPU, Trbvh may automatically fallback to build on the CPU.
Sbvh	The Split-BVH (SBVH) is a high quality bounding volume hierarchy. While build times are highest, it was traditionally the method of choice for static geometry due to its high ray tracing performance, but may be superseded by Trbvh. Improvements over regular BVHs are especially visible if the geometry is non-uniform (for example, in collections of triangles of different sizes). This builder can be used for any type of geometry, but for optimal performance with triangle geometry, specialized properties should be set (see Table 4)[3] .
Bvh	The Bvh builder constructs a classic bounding volume hierarchy. It has relatively good traversal performance and does not focus on fast construction performance, but it supports refitting for fast incremental updates (Table 4). Bvh is often the best choice for acceleration structures built over groups.
NoAccel	This is a dummy builder which does not construct an actual acceleration structure. Traversal loops over all elements and intersects each one with the ray. This is very inefficient for anything but very simple cases, but can sometimes outperform real acceleration structures, for example, on a group with very few child nodes.

Table

3 – Supported builders

Table 3 shows the builders currently available in OptiX. A builder is set using rtAccelerationSetBuilder. The builder can be changed at any time; switching builders will cause an acceleration structure to be flagged for rebuild.

This example shows a typical initialization of an acceleration object:

Listing 3.34

```
RTacceleration accel;
rtAccelerationCreate( context, &accel );
rtAccelerationSetBuilder( accel, "Trbvh" );
```

---

[2]See Tero Karras and Timo Aila, *Fast Parallel Construction of High-Quality Bounding Volume Hierarchies*. http://highperformancegraphics.org/wp-content/uploads/Karras-BVH.pdf
[3]See Martin Stich, Heiko Friedrich, Andreas Dietrich. *Spatial Splits in Bounding Volume Hierarchies*. http://www.nvidia.com/object/nvidia_research_pub_012.html

3.5.3   Acceleration structure properties

Fine-tuning acceleration structure construction can be useful depending on the situation. For this purpose, builders expose various named properties, which are listed in Table 4:

Property	Available	Description
`refit`	Bvh Trbvh	If set to 1, the builder will only readjust the node bounds of the bounding volume hierarchy instead of constructing it from scratch. Refit is only effective if there is an initial BVH already in place, and the underlying geometry has undergone relatively modest deformation. In this case, the builder delivers a very fast BVH update without sacrificing too much ray tracing performance. The default is 0.
`vertex_buffer_name`	Bvh Trbvh	The name of the buffer variable holding triangle vertex data. Each vertex consists of 3 floats. Optional for Sbvh (but recommended if the geometry consists of triangles). The default is `vertex_buffer`.
`vertex_buffer_stride`	Bvh Trbvh	The offset between two vertices in the vertex buffer, given in bytes. The default value is 0, which assumes the vertices are tightly packed.
`index_buffer_name`	Bvh Trbvh	The name of the buffer variable holding vertex index data. The entries in this buffer are indices of type `int`, where each index refers to one entry in the vertex buffer. A sequence of three indices represents one triangle. If no index buffer is given, the vertices in the vertex buffer are assumed to be a list of triangles, with every three vertices in a row forming a triangle. The default is `index_buffer`.
`index_buffer_stride`	Bvh Trbvh	The offset between two indices in the index buffer, given in bytes. The default value is 0, which assumes the indices are tightly packed.
`chunk_size`	Trbvh	Number of bytes to be used for a partitioned acceleration structure build. If no chunk size is set, or set to 0, the chunk size is chosen automatically. If set to −1, the chunk size is unlimited. The minimum chunk size is currently 64MB. Please note that specifying a small chunk size reduces the peak-memory footprint of the `Trbvh`, but can result in slower rendering performance.

Table 4 – Acceleration structure properties

Properties are specified using `rtAccelerationSetProperty`. Their values are given as strings, which are parsed by OptiX. Properties take effect only when an acceleration structure is actually rebuilt. Setting or changing the property does not itself mark the acceleration

structure for rebuild; see the next section for details on how to do that. Properties not recognized by a builder will be silently ignored.

Listing 3.35

```
rtAccelerationSetProperty( accel, "refit", "1" );
```
Enable fast refitting on a BVH acceleration.

3.5.4   Acceleration structure builds

In OptiX, acceleration structures are flagged (marked "dirty") when they need to be rebuilt. During `rtContextLaunch`, all flagged acceleration structures are built before ray tracing begins. Every newly created `RTacceleration` object is initially flagged dirty.

An application can decide at any time to explicitly mark an acceleration structure for rebuild. For example, if the underlying geometry of a geometry group changes, the acceleration structure attached to the geometry group must be recreated. This is achieved by calling `rtAccelerationMarkDirty`. This is also required if, for example, new child geometry instances are added to the geometry group, or if children are removed from it.

The same is true for acceleration structures on groups: adding or removing children, changing transforms below the group, etc., are operations which require the group's acceleration to be marked as dirty. As a rule of thumb, every operation that causes a modification to the underlying geometry over which the structure is built (in the case of a group, that geometry is the children's axis-aligned bounding boxes) requires a rebuild. However, no rebuild is required if, for example, some parts of the graph change further down the tree, without affecting the bounding boxes of the immediate children of the group.

Note that the application decides independently for each single acceleration structure in the graph whether a rebuild is necessary. OptiX will not attempt to automatically detect changes, and marking one acceleration structure as dirty will not propagate the dirty flag to any other acceleration structures. Failure to mark acceleration structures as dirty when necessary may result in unexpected behavior—usually missing intersections or performance degradation.

3.5.5   Shared acceleration structures

Mechanisms such as a graph node being attached as a child to multiple other graph nodes make composing the node graph flexible, and enable interesting instancing applications. Instancing can be seen as inexpensive reuse of scene objects or parts of the graph by referencing nodes multiple times instead of duplicating them.

OptiX decouples acceleration structures as separate objects from other graph nodes. Hence, acceleration structures can naturally be shared between several groups or geometry groups, as long as the underlying geometry on which the structure is built is the same:

Listing 3.36

```
rtGroupSetAcceleration( group1, accel );
rtGroupSetAcceleration( group2, accel );
rtGroupSetAcceleration( group3, accel );
```
Attach one acceleration to multiple groups.

Note that the application must ensure that each node sharing the acceleration structure has matching underlying geometry. Failure to do so will result in undefined behavior. Also, acceleration structures cannot be shared between groups and geometry groups.

The capability of sharing acceleration structures is a powerful concept to maximize efficiency, as shown in Figure 3.4. The acceleration node in the center of the figure is attached to both geometry groups, and both geometry groups reference the same geometry objects. This reuse of geometry and acceleration structure data minimizes both memory footprint and acceleration construction time. Additional geometry groups could be added in the same manner at very little overhead.
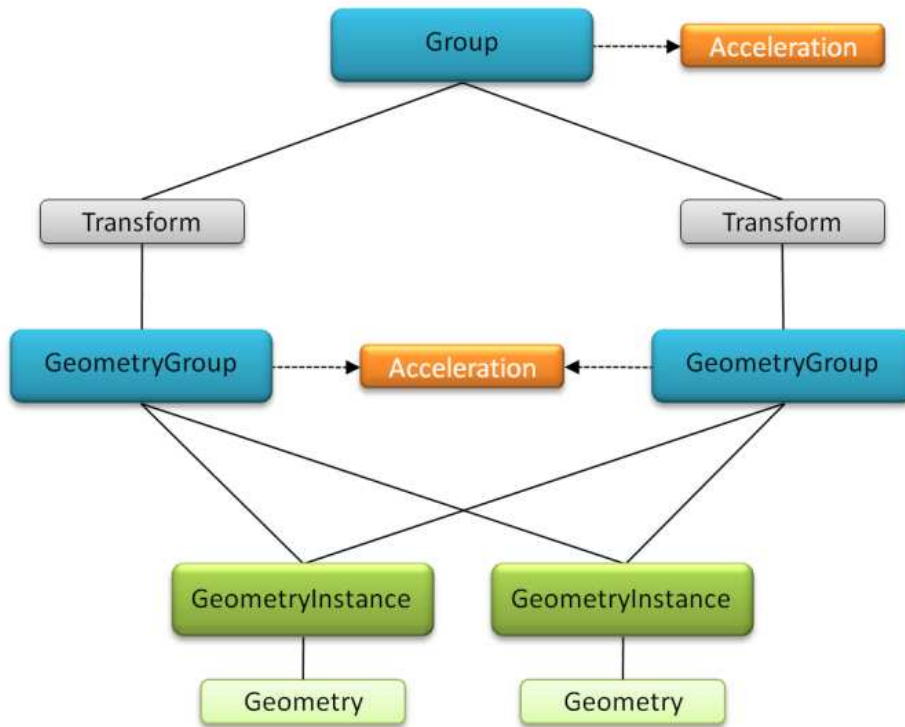
Fig. 3.4 – Two geometry groups sharing an acceleration structure and the underlying geometry objects.

3.6   Triangles

The `RTgeometrytriangles` type provides OptiX with built-in support for triangles. `RTgeometrytriangles` complements the `RTgeometry` type, with functions that can explicitly define the triangle data. Custom intersection and bounding box programs are not required by `RTgeometrytriangles`; the application only needs to provide the triangle data to OptiX. OptiX provides a built-in intersection program and a triangle-primitive-aware BVH-build mechanism. The `RTgeometrytriangles` triangle data required for acceleration structure traversal and intersection is directly stored within the acceleration structure.

3.6.1   `RTgeometrytriangles`

Two functions define the data of indexed and unorganized collections of triangles: `rtGeometryTrianglesSetIndexedTriangles` and `rtGeometryTrianglesSetTriangles`.[4]

---

[4]Non-indexed triangle sets are sometimes called "triangle soup" to emphasize the lack of a structural relationship between the triangles.

3.6.2   Defining unorganized triangle collections

Function `rtGeometryTrianglesSetTriangles` defines a buffer of vertices and how those
vertices should be interpreted as triangles.

Listing 3.37

```
RTresult rtGeometryTrianglesSetTriangles(
    RTgeometrytriangles  geometry,
    unsigned int         num_triangles,
    RTbuffer             vertex_buffer,
    RTsize               vertex_buffer_byte_offset,
    RTsize               vertex_byte_stride,
    RTformat             position_format,
    RTgeometrybuildflags build_flags)
```

The parameters of `rtGeometryTrianglesSetTriangles`:

`geometry`
GeometryTriangles node.

`num_triangles`
Number of triangles in `vertex_buffer`.

`vertex_buffer`
A set of triangles of `GeometryTriangles` geometry.

`vertex_buffer_byte_offset`
An optional offset to the first vertex in `vertex_buffer`.

`vertex_byte_stride`
The stride in bytes between vertices in `vertex_buffer`.

`position_format`
One of the following:

    RTformat::RT_FORMAT_FLOAT3

    RTformat::RT_FORMAT_HALF3

    RTformat::RT_FORMAT_FLOAT2

    RTformat::RT_FORMAT_HALF2

In case of `RTformat::RT_FORMAT_FLOAT2` and `RTformat::RT_FORMAT_HALF2` the third
component is assumed to be zero, which can be useful for planar geometry.

`build_flags`
If set to `RT_GEOMETRYBUILDFLAGS_RELEASE_BUFFERS`, the buffer holding the vertices will
be cleared once the acceleration structure is built and the triangle data has been copied to
it. Sharing the buffer with another GeometryTriangles node is possible if both geometries
are built within one OptiX launch. Setting `build_flags` to
`RT_GEOMETRYBUILDFLAGS_RELEASE_BUFFERS` will result in an error if refitting is defined
for the acceleration structure.

3.6.3   Defining indexed triangle collections

Listing 3.38 is the function signature of `rtGeometryTrianglesSetIndexedTriangles`:

Listing 3.38

```
RTresult rtGeometryTrianglesSetIndexedTriangles(
    RTgeometrytriangles  geometry,
    unsigned int         num_triangles,
    RTbuffer             index_buffer,
    RTsize               index_buffer_byte_offset,
    RTsize               tri_indices_byte_stride,
    RTformat             tri_indices_format,
    unsigned int         num_vertices,
    RTbuffer             vertex_buffer,
    RTsize               vertex_buffer_byte_offset,
    RTsize               vertex_byte_stride,
    RTformat             position_format,
    RTgeometrybuildflags build_flags );
```

Like `rtGeometryTrianglesSetTriangles`, function
`rtGeometryTrianglesSetIndexedTriangles` defines the following parameters:

> `geometry`
>
> `num_triangles`
>
> `vertex_buffer`
>
> `vertex_buffer_byte_offset`
>
> `vertex_byte_stride`
>
> `position_format`
>
> `build_flags`

For the definition of indexed triangles, the `rtGeometryTrianglesSetIndexedTriangles` adds
these parameters:

`index_buffer`
> Indices of the triangle vertices in `vertex_buffer`. The indices are grouped as triplets to
> define triangles. The number of indices will be three times the value of `num_triangles`.

`index_buffer_byte_offset`
> An optional offset to the first index in `index_buffer`.

`tri_indices_byte_stride`
> Sets the stride in bytes between triplets of indices. Spacing is only supported between
> triplets; no spacing is allowed between indices within a triplet.

`tri_indices_format`
> One of two values:
>
> > `RTformat::RT_FORMAT_UNSIGNED_INT3`
> >
> > `RTformat::RT_FORMAT_UNSIGNED_SHORT3`

`num_vertices`
> Defines the number of vertices stored in `vertex_buffer`.

3.6.4 Additional `RTgeometrytriangles` functions

Creation, destruction, validation and inspection of `RTgeometrytriangles` instances are provided by six additional functions:

Listing 3.39

```
RTresult rtGeometryTrianglesCreate(
    RTcontext context, RTgeometrytriangles* geometry)

RTresult rtGeometryTrianglesDestroy(
    RTgeometrytriangles geometry)

RTresult rtGeometryTrianglesValidate(
    RTgeometrytriangles geometry)

RTresult rtGeometryTrianglesGetContext(
    RTgeometrytriangles geometry, RTcontext* context)

RTresult rtGeometryTrianglesGetPrimitiveIndexOffset(
    RTgeometrytriangles geometrytriangles, unsigned int* index_offset)

RTresult rtGeometryTrianglesSetPrimitiveIndexOffset(
    RTgeometrytriangles geometrytriangles, unsigned int index_offset)
```

3.6.5 The `RTgeometrybuildflags` enum

The `RTgeometrybuildflags` enum can be used to set a combination of build flags. It consists of one option:

Listing 3.40

```
typedef enum {
    RT_GEOMETRYBUILDFLAGS_NONE = 0x00,    No special flags set

    RT_GEOMETRYBUILDFLAGS_RELEASE_BUFFERS = 0x01   User buffers are released after
                                                    consumption
} RTgeometrybuildflags;
```

3.6.6 Differences between `RTgeometrytriangles` and `RTgeometry`

Note that in comparison to `RTgeometry`, a number of analogous functions do *not* exist for `RTgeometrytriangles`:

    setPrimitiveCount

    getPrimitiveCount

    setIntersectionProgram

    getIntersectionProgram

    setBoundingBoxProgram

```
getBoundingBoxProgram
```

In addition, `RTgeometrytriangles` does not support variables.

The primitive count is set in either of the two functions above as parameter `num_triangles`. An offset can be set for each buffer. Since `RTgeometrytriangles` comes with a built-in intersection program, a custom program cannot be set. Variables cannot be set as the new type does not provide a valid scope to any other program type (such as any-hit (page 53), closest-hit (page 52), miss (page 55), etc.)

3.6.7 Adding `RTgeometrytriangles` to a scene graph

A geometry instance of type `RTgeometrytriangles` can be attached to a GeometryInstances just like a `RTgeometry`.

Listing 3.41

```
rtGeometryInstanceSetGeometryTriangles(
    RTgeometryinstance instance, RTgeometrytriangles geometry)


rtGeometryInstanceGetGeometryTriangles(
    RTgeometryinstance instance, RTgeometrytriangles* geometry)
```

Only one material can be set on the geometry instance if a `RTgeometrytriangles` is attached. This material is used for any triangle intersection with the `RTgeometrytriangles`.

3.6.8 Remote launches

The `rtContextLaunch` calls offered by traditional OptiX work for remote rendering but are not well suited to it. Because the API calls are synchronous, each launch/display cycle incurs the full network round-trip latency, and thus performance is usually not acceptable over standard network connections. Progressive launches, on the other hand, are asynchronous, and can achieve smooth application performance even over a high latency connection.

3.6.9 Parallelization

OptiX can parallelize work across a small number of local GPUs. Progressive renderers, such as path tracers, are one of the most common use cases of OptiX. The fact that the image samples they compute are independent of each other provides a natural way to parallelize the problem. The Progressive Launch API therefore combines the assumption that work can be split into many independent parts with the capability to launch kernels asynchronously.

3.6.10 Progressive launches

While most existing OptiX applications will work unchanged when run on a remote device (see "Limitations" (page 39) for caveats), progressive launches must be used instead of `rtContextLaunch` for optimal performance.

Instead of requesting the generation of a single frame, a progressive launch, triggered by `rtContextLaunchProgressive2D`, requests *multiple subframes* at once. A subframe is output buffer content which is composited with other subframes to yield the final frame. In most progressive renderers, this means that a subframe simply contains a single sample per pixel.

Progressive launch calls are non-blocking. An application typically executes a progressive launch, and then continuously polls the *stream buffers* associated with its output, using `rtBufferGetProgressiveUpdateReady`. If that call reports that an update is available, the stream buffer can be mapped and the content displayed.

If any OptiX API functions are called while a progressive launch is in progress, the launch will stop generating subframes until the next time a progressive launch is triggered (the exception is the API calls to poll and map the stream buffers). This way, state changes to OptiX, such as setting variables using `rtVariableSet`, can be made easily and efficiently in combination with a render loop that polls for stream updates and executes a progressive launch. This method is outlined in the example pseudocode below.

3.6.11   Stream buffers

Accessing the results of a progressive launch is typically done through a new type of buffer called a *stream buffer*. Stream buffers allow the remotely rendered frames to be sent to the application client using compressed video streaming, greatly improving response times while still allowing the application to use the result frame in the same way as with a non-stream output buffer.

Stream buffers are created using `rtBufferCreate` with the type set to `RT_BUFFER_PROGRESSIVE_STREAM`. A stream buffer must be bound to a regular output buffer via `rtBufferBindProgressiveStream` in order to define its data source.

By executing the bind operation, the system enables automatic compositing for the output buffer. That is, any values written to the output buffer by device code will be averaged into the stream buffer, rather than overwriting the previous value as in regular output buffers. Compositing happens automatically and potentially in parallel across many devices on the network, or locally if remote rendering is not used.

Several configuration options are available for stream buffers, such as the video stream format to use, and parameters to trade off quality versus speed. Those options can be set using `rtBufferSetAttribute`. Note that some of the options only take effect if rendering happens on a remote device, and are a no-op when rendering locally. This is because stream buffers don't undergo video compression when they don't have to be sent across a network.

In addition to automatic compositing, the system also tone-maps and quantizes the averaged output before writing it into a stream. Tone mapping is performed using a simple built-in operator with a user- defined gamma value (specified using `rtBufferSetAttribute`). The tone-map operator is defined as:

Listing 3.42
`final_value = clamp( pow( hdr_value, 1/gamma ), 0, 1 )`

Accessing a progressive stream happens by mapping the stream buffer, just like any regular buffer, and reading out the frame data. The data is uncompressed, if necessary, when mapped. The data available for reading will always represent the most recent update to the stream if a progressive launch is in progress, so a frame that is not read on time may be skipped (for example, if polling happens at a low frequency).

It is also possible to map an output buffer that is bound as a data source for a stream buffer. This can be useful to access "final frame" data—the uncompressed and unquantized accumulated output. Note that mapping a non-stream buffer will cause the progressive launch to stop generating subframes, and that such a map operation is much slower than mapping a stream.

3.6.12   Device code

In OptiX device code, the subframe index used for progressive rendering is exposed as a semantic variable of type `unsigned int`. Its value is guaranteed to be unique for each subframe in the current progressive launch, starting at zero for the first subframe and increasing by one with each subsequent subframe. For example, an application performing stochastic sampling may use this variable to seed a random number generator. (For a description of semantic variables, see "Communication through variables" (page 42).)

The current subframe index can be accessed in shader programs by declaring the following variable:

Listing 3.43

```
rtDeclareVariable(unsigned int, index, rtSubframeIndex,);
```

Computed pixel values can be written to an output buffer, just like for non-progressive rendering. Output buffers that are bound as sources to stream buffers will then be averaged automatically and processed (as described in "Buffers" (page 15)).

Note in particular that device code does not use stream buffers directly.

3.6.13   Limitations

- OpenGL interoperability is limited with remote rendering, and using it is discouraged for performance reasons. See "Interoperability with OpenGL" (page 89) for more information. Direct3D and CUDA interop are not supported.

- Using buffers of type `RT_BUFFER_INPUT_OUTPUT` in combination with remote rendering yields undefined results.

- `rtPrintf` and associated host functions are not supported in combination with remote rendering.

- `rtContextSetTimeoutCallback` is not supported in combination with remote rendering.

- Output buffers used as data sources for progressive stream buffers must be of `RT_FORMAT_FLOAT3` or `RT_FORMAT_FLOAT4` format. For performance reasons, using `RT_FORMAT_FLOAT4` is strongly recommended.

- Stream buffers must be of `RT_FORMAT_UNSIGNED_BYTE4` format.

---

4 Programs

This chapter describes the different kinds of OptiX programs, which provide programmatic control over ray intersection, shading, and other general computation in OptiX ray tracing kernels. OptiX programs are associated with binding points serving different semantic roles during a ray tracing computation. Like other concepts, OptiX abstracts programs through its object model as *program objects*.

4.1 OptiX program objects

The central theme of the OptiX API is programmability. OptiX programs are written in CUDA C, and specified to the API through a string or file containing PTX, the parallel thread execution virtual assembly language associated with CUDA. The nvcc compiler that is distributed with the CUDA SDK is used to create PTX in conjunction with the OptiX header files.

These PTX files are then bound to Program objects via the host API. Program objects can be used for any of the OptiX program types discussed later in this section.

4.1.1 Managing program objects

OptiX provides two API entry points for creating `Program` objects: `rtProgramCreateFromPTXString`, and `rtProgramCreateFromPTXFile`. The former creates a new Program object from a string of PTX source code. The latter creates a new Program object from a file of PTX source on disk:

Listing 4.1

```
RTcontext context = ...;
const char *ptx_filename = ...;
const char *program_name = ...;
RTprogram program = ...;
rtProgramCreateFromPTXFile(
    context, ptx_filename, function_name, &program );
```

In this example, `ptx_filename` names a file of PTX source on disk, and `function_name` names a particular function of interest within that source. If the program is ill-formed and cannot compile, these entry points return an error code.

Program objects may be checked for completeness using the `rtProgramValidate` function, as the following example demonstrates:

Listing 4.2

```
if ( rtProgramValidate(context, program)!=RT_SUCCESS ) {
    printf( "Program is not complete." );
}
```

An error code returned from `rtProgramValidate` indicates an error condition due to the program object or any other objects bound to it.

Finally, the `rtProgramGetContext` function reports the context object owning the program object, while `rtProgramDestroy` invalidates the object and frees all resources related to it.

4.1.2 Communication through variables

OptiX program objects communicate with the host program through variables. Variables are declared in an OptiX program using the `rtDeclareVariable` macro:

Listing 4.3

```
rtDeclareVariable( float, x, , );
```

This declaration creates a variable named *x* of type `float` which is available to both the host program through the OptiX variable object API, and to the device program code through usual C language semantics. Notice that the last two arguments are left blank in this example. The commas must still be specified.

Taking the address of a variable on the device is not supported. This means that pointers and references to x in the above example are not allowed. If, for instance, you needed to pass x into a function taking a `float*` argument you would need to first copy x into a stack variable and then pass in the address of this local variable:

Listing 4.4

```
void my_func( float* my_float) {...}

RT_PROGRAM call_my_func()
{
    my_func(&x);         Not allowed
    float local_x = x;
    my_func(&local_x);    Allowed
}
```

Variables declared in this way may be read and written by the host program through the `rtVariableGet` and `rtVariableSet` family of functions. When variables are declared this way, they are implicitly const-qualified from the device program's perspective. If communication from the program to the host is necessary, an `rtBuffer` should be used instead.

As of OptiX 2.0, variables may be declared inside arbitrarily nested namespaces to avoid name conflicts. References from the host program to namespace-enclosed OptiX variables will need to include the full namespace.

Program variables may also be declared with *semantics*. Declaring a variable with a semantic binds the variable to a special value which OptiX manages internally over the lifetime of the ray tracing kernel. For example, declaring a variable with the `rtCurrentRay` semantic creates a special read-only program variable that mirrors the value of the `Ray` currently being traced through the program flow:

Listing 4.5

```
rtDeclareVariable( Ray, ray, rtCurrentRay, );
```

Variables declared with a built-in semantic exist only during ray tracing kernel run time and may not be modified or queried by the host program. Unlike regular variables, some semantic variables may be modified by the device program.

Declaring a variable with an *annotation* associates with it a read-only string which, for example, may be interpreted by the host program as a human-readable description of the variable. For example:

Listing 4.6

```
rtDeclareVariable( float, shininess, , "The shininess of the sphere" );
```

A variable's annotation is the fourth argument of `rtDeclareVariable`, following the variable's optional semantic argument. The host program may query a variable's annotation with the `rtVariableGetAnnotation` function.

4.1.3 Internally provided semantics

OptiX manages five internal semantics for program variable binding. Table 5 summarizes in which types of program these semantics are available, along with their access rules from device programs and a brief description of their meaning.

Semantic name	Access	Description	Ray generation	Exception	Closest hit	Any hit	Miss	Intersection	Bounding box	Visit
`rtLaunchIndex`	read only	The unique index identifying each thread launched by `rtContextLaunch`*n*D.	✓	✓	✓	✓	✓	✓		✓
`rtCurrentRay`	read only	The state of the current ray.			✓	✓	✓	✓		✓
`rtPayload`	read/write	The state of the current ray's payload of user-defined data.			✓	✓	✓			✓
`rtIntersectionDistance`	read only	The parametric distance from the current ray's origin to the closest intersection point yet discovered.			✓	✓		✓		✓
`rtSubframeIndex`	read only	The unique index identifying each subframe in a progressive launch. Zero for non-progressive launches.	✓	✓	✓	✓	✓	✓		✓

Table 5 – Semantic variables

4.1.4 Attribute variables

In addition to the semantics provided by OptiX, variables may also be declared with user-defined semantics called *attributes*. Unlike built-in semantics, the value of variables

---

declared in this way must be managed by the programmer. Attribute variables provide a mechanism for communicating data between the intersection program and the shading programs (for example, surface normals and texture coordinates). Attribute variables may *only* be written in an intersection program between calls to `rtPotentialIntersection` and `rtReportIntersection`. Although OptiX may not find all object intersections in order along the ray, the value of the attribute variable is guaranteed to reflect the value at the closest intersection at the time that the closest-hit program is invoked.

> **Note:** Because intersections may not be found in order, programs should use attribute variables (as opposed to the ray payload) to communicate information about the local hit point between intersection and shading programs.

The following example declares an attribute variable of type `float3` named *normal*. The semantic association of the attribute is specified with the user-defined name `normal_vec`. This name is arbitrary, and is the link between the variable declared here and another variable declared in the closest-hit program. The two attribute variables need not have the same name as long as their attribute names match.

Listing 4.7

```
rtDeclareVariable( float3, normal, attribute normal_vec, );
```

4.1.5   Program variable scoping

OptiX program variables can have their values defined in two ways: static initializations, and (more typically) by variable declarations attached to API objects. A variable declared with a static initializer will only use that value if it does not find a definition attached to an API object. A declaration with static initialization is written:

Listing 4.8

```
rtDeclareVariable( float, x, , ) = 5.0f;
```

The OptiX variable scoping rules provide a valuable inheritance mechanism that is designed to create compact representations of material and object parameters. To enable this, each program type also has an ordered list of scopes through which it will search for variable definitions in order. For example, a closest-hit program that refers to a variable named `color` will search the `Program`, `GeometryInstance`, `Material` and `Context` API objects for definitions created with the `rt*DeclareVariable` functions, in that order. Similar to scoping rules in a programming language, variables in one scope will shadow those in another scope. summarizes the scopes that are searched for variable declarations for each type of program.

Program type	Search order			
Closest hit Any hit	`Program`	`GeometryInstance`	`Material`	`Context`
Intersection Bounding box	`Program`	`GeometryInstance`	`Geometry`	`Context`
Ray generation Exception Miss	`Program`	`Context`		
Visit	`Program`	`Node`		

Table 6 – Scope search order for each type of program (from left to right)

It is possible for a program to find multiple definitions for a variable in its scopes depending upon where the program is called. For example, a closest-hit program may be attached to several `Material` objects and reference a variable named `shininess`. We can attach a variable definition to the `Material` object as well as attach a variable definition to specific GeometryInstance objects that we create that reference that `Material`.

During execution of a specific GeometryInstance's closest-hit program, the value of *shininess* depends on whether the particular instance has a definition attached: if the GeometryInstance defines shininess, then that value will be used. Otherwise, the value will be taken from the `Material` object. As you can see from Table 6 above, the program searches the GeometryInstance scope before the `Material` scope. Variables with definitions in multiple scopes are said to be *dynamic* and may incur a performance penalty. Dynamic variables are therefore best used sparingly.

4.1.6   Program variable transformation

Recall that rays have a projective transformation applied to them upon encountering `Transform` nodes during traversal. The transformed ray is said to be in *object space*, while the original ray is said to be in *world space*.

Programs with access to the `rtCurrentRay` semantic operate in the spaces summarized in Table 7:

Program type	Space
Closest hit	world
Any hit	object
Miss	world
Intersection	object
Visit	object

Table 7 – Ray space

To facilitate transforming variables from one space to another, OptiX's CUDA C API provides a set of functions:

Listing 4.9

```
__device__ float3 rtTransformPoint(
               RTtransformkind kind, const float3& p )
__device__ float3 rtTransformVector(
```

```
                    RTtransformkind kind, const float3& v )
 __device__ float3 rtTransformNormal(
                    RTtransformkind kind, const float3& n )
 __device__ void rtGetTransform(
                    RTtransformkind kind, float matrix[16] )
```

The first three functions transform a `float3`—interpreted as a point, vector, or normal vector—from object space to world space or vice versa, depending on the value of a `RTtransformkind` flag passed as an argument. `rtGetTransform` returns the four-by-four matrix representing the current transformation from object to world space (or vice versa, depending on the `RTtransformkind` argument). For best performance, use the CUDA C `rtTransform*` functions rather than performing your own explicit matrix multiplication with the result of `rtGetTransform`.

A common use case of variable transformation occurs when interpreting attributes passed from the intersection program to the closest-hit program. Intersection programs often produce attributes, such as normal vectors, in object space. Should a closest-hit program wish to consume that attribute, it often must transform the attribute from object space to world space:

Listing 4.10

```
float3 n = rtTransformNormal( RT_OBJECT_TO_WORLD, normal );
```

4.2   The program scope of API function calls

Not all OptiX function calls are supported in all types of user-provided programs. For example, it doesn't make sense to spawn a new ray inside an intersection program, so this behavior is disallowed. A complete table of what device-side functions are allowed is given below. (See "Callable programs" (page 60).)

	Ray generation	Exception	Closest hit	Any hit	Miss	Intersection	Bounding box	Visit	Bindless callable program
`rtTransform*`			✓	✓	✓	✓	✓	✓	
`rtTrace`	✓		✓		✓				
`rtThrow`	✓		✓	✓	✓	✓	✓	✓	✓
`rtPrintf`	✓	✓	✓	✓	✓	✓	✓	✓	✓
`rtTerminateRay`				✓					
`rtIgnoreIntersection`				✓					
`rtIntersectChild`								✓	
`rtPotentialIntersection`						✓			
`rtReportIntersection`						✓			
Callable program (page 60)	✓	✓	✓	✓	✓	✓	✓	✓	✓

Table 8 – Scopes allowed for device-side API functions

4.3   Ray generation programs

A *ray generation program* serves as the first point of entry upon a call to `rtContextLaunch`*n*D. As such, it serves a role analogous to the main function of a C program. Like C's `main` function, any subsequent computation performed by the kernel, from casting rays to reading and writing from buffers, is spawned by the ray generation program. However, unlike a serial C program, an OptiX ray generation program is executed many times in parallel—once for each thread implied by `rtContextLaunch`*n*D's parameters.

Each thread is assigned a unique `rtLaunchIndex`. The value of this variable may be used to distinguish the thread from its neighbors, for example, to enable writing to a unique location in an `rtBuffer`:

Listing 4.11

```
rtBuffer<float, 1> output_buffer;
rtDeclareVariable( unsigned int, index, rtLaunchIndex, );
...;
float result = ...;
output_buffer[index] = result;
```

In this case, the result is written to a unique location in the output buffer. In general, a ray generation program may write to any location in output buffers, as long as care is taken to avoid race conditions between buffer writes.

4.3.1   Entry point indices

To configure a ray tracing kernel launch, the programmer must specify the desired ray generation program using an *entry point index*. The total number of entry points for a context is specified with `rtContextSetEntryPointCount`:

Listing 4.12

```
RTcontext context = ...;
unsigned int num_entry_points = ...;
rtContextSetEntryPointCount( context, num_entry_points );
```

OptiX requires that each entry point index created in this manner have a ray generation program associated with it. A ray generation program may be associated with multiple indices. Use the `rtContextSetRayGenerationProgram` function to associate a ray generation program with an entry point index in the range [0, `num_entry_points`).

Listing 4.13

```
RTprogram prog = ...;
unsigned int index = ...;    Value of index is >= 0 and < num_entry_points
rtContextSetRayGenerationProgram( context, index, prog );
```

4.3.2   Launching a ray generation program

`rtContextLaunch`*n*D takes as a parameter the entry point index of the ray generation program to launch:

Listing 4.14

```
RTsize width = ...;
rtContextLaunch1D( context, index, width );
```

If no ray generation program has been associated with the entry point index specified by `rtContextLaunch`*n*D's parameter, the launch will fail.

4.3.3   Ray generation program function signature

In CUDA C, ray generation programs return `void` and take no parameters. Like all OptiX programs, ray generation programs written in CUDA C must be tagged with the `RT_PROGRAM` qualifier. The following snippet shows an example ray generation program function prototype:

Listing 4.15

```
RT_PROGRAM void ray_generation_program( void );
```

4.3.4   Example ray generation program

The following example ray generation program implements a pinhole camera model in a rendering application. This example demonstrates that ray generation programs act as the gateway to all ray tracing computation by initiating traversal through the `rtTrace` function, and often store the result of a ray tracing computation to an output buffer.

Note the variables `eye`, `U`, `V`, and `W`. Together, these four variables allow the host API to specify the position and orientation of the camera.

Listing 4.16

```
rtBuffer<uchar4, 2> output_buffer;
rtDeclareVariable( uint2, index, rtLaunchIndex, );
rtDeclareVariable( rtObject, top_object, , );

rtDeclareVariable(float3, eye, , );
rtDeclareVariable(float3, U, , );          Define camera position and orientation
rtDeclareVariable(float3, V, , );
rtDeclareVariable(float3, W, , );

struct Payload {
    uchar4 result;
};

RT_PROGRAM void pinhole_camera( void )
{
    uint2 screen = output_buffer.size();

    float2 d = (make_float2( index ) / make_float2( screen )) * 2.f - 1.f;
    float3 origin = eye;
    float3 direction = normalize( d.x*U + d.y*V + W );

    Ray ray = make_Ray(
        origin, direction, 0, 0.05f, RT_DEFAULT_MAX );    Define ray from camera
                                                           position and direction

    Payload payload;
    rtTrace( top_object, ray, payload );
```

```
    output_buffer[index] = payload.result;
}
```

4.4 Exception programs

OptiX ray tracing kernels invoke an *exception program* when certain types of serious errors are encountered. Exception programs provide a means of communicating to the host program that something has gone wrong during a launch. The information an exception program provides may be useful in avoiding an error state in a future launch or for debugging during application development.

4.4.1 Exception program entry point association

An exception program is associated with an entry point using the `rtContextSetExceptionProgram` function:

Listing 4.17

```
RTcontext context = ...;
RTprogram program = ...;
unsigned int index = ...;      Value of index is >= 0 and < num_entry_points
rtContextSetExceptionProgram( context, index, program );
```

Unlike with ray generation programs, the programmer need not associate an exception program with an entry point. By default, entry points are associated with an internally provided exception program that silently ignores errors.

As with ray generation programs, a single exception program may be associated with many different entry points.

4.4.2 Exception types

OptiX detects a number of different error conditions that result in exception programs being invoked. An exception is identified by its code, which is an integer defined by the OptiX API. For example, the exception code for the stack overflow exception is `RT_EXCEPTION_STACK_OVERFLOW`.

The type or code of a caught exception can be queried by calling `rtGetExceptionCode` from the exception program. More detailed information on the exception can be printed to the standard output using `rtPrintExceptionDetails`.

In addition to the built in exception types, OptiX provides means to introduce user-defined exceptions. Exception codes between `RT_EXCEPTION_USER` (0x400) and `0xFFFF` are reserved for user exceptions. To trigger such an exception, `rtThrow` is used:

Listing 4.18

```
#define MY_EXCEPTION_0 RT_EXCEPTION_USER + 0
#define MY_EXCEPTION_1 RT_EXCEPTION_USER + 1      Define user-specified exception codes.

RT_PROGRAM void some_program()
```

```
{
    ...
    if ( condition0 )
        rtThrow( MY_EXCEPTION_0 );        Throw user exceptions from within a program.
    if ( condition1 )
        rtThrow( MY_EXCEPTION_1 );
    ...
}
```

In order to control the run-time overhead involved in checking for error conditions, individual types of exceptions may be switched on or off using `rtContextSetExceptionEnabled`. Disabling exceptions usually results in faster performance, but is less safe. By default, only `RT_EXCEPTION_STACK_OVERFLOW` is enabled. During debugging, it is often useful to turn on all available exceptions. This can be achieved with a single call:

Listing 4.19

```
...
rtContextSetExceptionEnabled(context, RT_EXCEPTION_ALL, 1);
...
```

4.4.3   Exception program function signature

In CUDA C, exception programs return `void`, take no parameters, and use the `RT_PROGRAM` qualifier:

Listing 4.20

```
RT_PROGRAM void exception_program( void );
```

4.4.4   Example exception program

The following example code demonstrates a simple exception program which indicates a stack overflow error by outputting a special value to an output buffer which is otherwise used as a buffer of pixels. In this way, the exception program indicates the `rtLaunchIndex` of the failed thread by marking its location in a buffer of pixels with a known color. Exceptions which are not caused by a stack overflow are reported by printing their details to the console.

Listing 4.21

```
rtDeclareVariable( int, launch_index, rtLaunchIndex, );
rtDeclareVariable( float3, error, , ) = make_float3(1,0,0);
rtBuffer<float3, 2> output_buffer;

RT_PROGRAM void exception_program( void )
{
    const unsigned int code = rtGetExceptionCode();

    if ( code == RT_EXCEPTION_STACK_OVERFLOW )
```

```
        output_buffer[launch_index] = error;
    else
        rtPrintExceptionDetails();
}
```

4.5 Closest-hit programs

After a call to the `rtTrace` function, OptiX invokes a *closest-hit program* once it identifies the nearest primitive intersected along the ray from its origin. Closest-hit programs are useful for performing primitive-dependent processing that should occur once a ray's visibility has been established. A closest-hit program may communicate the results of its computation by modifying per-ray data or writing to an output buffer. It may also recursively call the `rtTrace` function. For example, a computer graphics application might implement a surface shading algorithm with a closest hit program.

4.5.1 Closest-hit program material association

A closest-hit program is associated with each (*material*, *ray_type*) pair. Each pair's default program is a no-op. This is convenient when an OptiX application requires many types of rays but only a small number of those types require special closest-hit processing.

The programmer may change an association with the `rtMaterialSetClosestHitProgram` function:

Listing 4.22

```
RTmaterial material = ...;
RTprogram program = ...;
unsigned int type = ...;
rtMaterialSetClosestHitProgram( material, type, program );
```

4.5.2 Closest-hit program function signature

In CUDA C, closest-hit programs return `void`, take no parameters, and use the `RT_PROGRAM` qualifier:

Listing 4.23

```
RT_PROGRAM void closest_hit_program( void );
```

4.5.3   Recursion in a closest-hit program

Though the `rtTrace` function is available to all programs with access to the `rtLaunchIndex` semantic, a common use case of closest-hit programs is to perform recursion by tracing more rays upon identification of the closest surface intersected by a ray. For example, a computer graphics application might implement Whitted-style ray tracing by recursive invocation of `rtTrace` and closest-hit programs. Care must be used to limit the recursion depth to avoid stack overflow.

4.5.4   Example closest-hit program

The following code example demonstrates a closest-hit program that transforms the normal vector computed by an intersection program (not shown) from the intersected primitive's local coordinate system to a global coordinate system. The transformed normal vector is returned to the calling function through a variable declared with the `rtPayload` semantic. Note that this program is quite trivial; normally the transformed normal vector would be used by the closest hit program to perform some calculation, for example, lighting. See the introductory tutorials[5] in the OptiX documentation set for examples.

Listing 4.24

```
rtDeclareVariable( float3, normal, attribute normal_vec, );
struct Payload {
    float3 result;
};
rtDeclareVariable( Payload, ray_data,rtPayload, );

RT_PROGRAM void closest_hit_program( void )
{
    float3 norm;
    norm = rtTransformNormal( RT_OBJECT_TO_WORLD, normal );
    norm = normalize( norm );
    ray_data.result = norm;
}
```

4.6   Any-hit programs

Instead of the closest intersected primitive, an application may wish to perform some computation for *any* primitive intersection that occurs along a ray cast during the `rtTrace` function; this usage model can be implemented using *any-hit programs*. For example, a rendering application may require some value to be accumulated along a ray at each surface intersection.

---

[5]`http://raytracing-docs.nvidia.com/optix/tutorials/index.html`

4.6.1   Any-hit program material association

Like closest-hit programs, an any-hit program is associated with each (*material*, *ray_type*) pair. Each pair's default association is with an internally provided any-hit program which implements a no-op.

The `rtMaterialSetAnyHitProgram` function changes the association of a (`material`, `ray_type`) pair:

Listing 4.25

```
RTmaterial material = ...;
RTprogram program = ...;
unsigned int type = ...;
rtMaterialSetAnyHitProgram( material, type, program );
```

4.6.2   Termination in an any-hit program

A common OptiX usage pattern is for an any-hit program to halt ray traversal upon discovery of an intersection. The any-hit program can do this by calling `rtTerminateRay`. This technique can increase performance by eliminating redundant traversal computations when an application only needs to determine whether any intersection occurs and identification of the nearest intersection is irrelevant. For example, a rendering application might use this technique to implement shadow ray casting, which is often a binary true or false computation.

4.6.3   Any-hit program function signature

In CUDA C, any-hit programs return `void`, take no parameters, and use the `RT_PROGRAM` qualifier:

Listing 4.26

```
RT_PROGRAM void any_hit_program( void );
```

4.6.4   Example any-hit program

The following code example demonstrates an any-hit program that implements early termination of shadow ray traversal upon intersection. The program also sets the value of a per-ray payload member, `attenuation`, to zero to indicate the material associated with the program is totally opaque.

Listing 4.27

```
struct Payload {
    float attenuation;
};

rtDeclareVariable( Payload, payload,rtPayload, );

RT_PROGRAM void any_hit_program( void )
{
```

```
    payload.attenuation = 0.f;
    rtTerminateRay();
}
```

4.7   Miss programs

When a ray traced by the `rtTrace` function intersects no primitive, a *miss program* is invoked. Miss programs may access variables declared with the `rtPayload` semantic in the same way as closest-hit and any-hit programs.

4.7.1   Miss program function signature

In CUDA C, miss programs return `void`, take no parameters, and use the `RT_PROGRAM` qualifier:

Listing 4.28

```
RT_PROGRAM void miss_program( void );
```

4.7.2   Example miss program

In a computer graphics application, the miss program may implement an environment mapping algorithm using a simple gradient, as this example demonstrates:

Listing 4.29

```
rtDeclareVariable( float3, environment_light, , );
rtDeclareVariable( float3, environment_dark, , );
rtDeclareVariable( float3, up, , );

struct Payload {
    float3 result;
};

rtDeclareVariable( Payload, payload,rtPayload, );
rtDeclareVariable( Ray, ray, rtCurrentRay, );

RT_PROGRAM void miss(void)
{
    float t = max( dot( ray.direction, up ), 0.0f );
    payload.result = lerp( environment_light, environment_dark, t );
}
```

4.8   Intersection and bounding box programs

Intersection and *bounding box programs* represents geometry by implementing ray-primitive intersection and bounding algorithms. These program types are associated with and queried from Geometry objects using `rtGeometrySetIntersectionProgram`, `rtGeometryGetIntersectionProgram`, `rtGeometrySetBoundingBoxProgram`, and `rtGeometryGetBoundingBoxProgram`.

4.8.1   Intersection and bounding box program function signatures

Like the previously discussed OptiX programs, in CUDA C, intersection and bounding box programs return `void` and use the `RT_PROGRAM` qualifier. Because Geometry objects are collections of primitives, these functions require a parameter to specify the index of the primitive of interest to the computation. This parameter is always in the range $[0, N)$, where N is given by the argument to the `rtGeometrySetPrimitiveCount` function.

Additionally, the bounding box program requires an array of floats to store the result of the bounding box computation, yielding these function signatures:

Listing 4.30

```
RT_PROGRAM void intersection_program( int prim_index);
RT_PROGRAM void bounding_box_program( int prim_index, float result[6]);
```

4.8.2 Reporting intersections

Ray traversal invokes an intersection program when the current ray encounters one of a Geometry object's primitives. It is the responsibility of an intersection program to compute whether the ray intersects with the primitive, and to report the parametric *t-value* of the intersection. Additionally, the intersection program is responsible for computing and reporting any details of the intersection, such as surface normal vectors, through attribute variables.

Once the intersection program has determined the t-value of a ray- primitive intersection, it must report the result by calling a pair of OptiX functions, `rtPotentialIntersection` and `rtReportIntersection`:

Listing 4.31

```
__device__ bool rtPotentialIntersection( float tmin )
__device__ bool rtReportIntersection( unsigned int material )
```

`rtPotentialIntersection` takes the intersection's t-value as an argument. If the t-value could potentially be the closest intersection of the current traversal the function narrows the *t-interval* of the current ray accordingly and returns `true`. If the t-value lies outside the t-interval the function returns `false`, whereupon the intersection program may trivially return.

If `rtPotentialIntersection` returns `true`, the intersection program may then set any attribute variable values and must subsequently call `rtReportIntersection`. This function takes an `unsigned int` specifying the index of a material that must be associated with an any-hit and closest-hit program. This material index can be used to support primitives of several different materials flattened into a single Geometry object. Traversal then immediately invokes the corresponding any-hit program. Should that any-hit program invalidate the intersection via the `rtIgnoreIntersection` function, then `rtReportIntersection` will return `false`. Otherwise, it will return `true`.

The values of attribute variables must be modified only between the call to `rtPotentialIntersection` and the call to `rtReportIntersection`. The result of writing to an attribute variable outside the bounds of these two calls is undefined. The values of attribute variables written in this way are accessible by any-hit and closest-hit programs.

If the any-hit program invokes `rtIgnoreIntersection`, any attributes computed will be reset to their previous values and the previous t-interval will be restored.

If no intersection exists between the current ray and the primitive, an intersection program need only return.

4.8.3 Specifying bounding boxes

Acceleration structures use *bounding boxes* to bound the spatial extent of scene primitives to accelerate the performance of ray traversal. A bounding box program's responsibility is to describe the minimal three dimensional axis-aligned bounding box that contains the primitive specified by its first argument and store the result in its second argument. Bounding boxes are always specified in object space, so the user should not apply any transformations to them.

For correct results bounding boxes must merely contain the primitive. For best performance bounding boxes should be as tight as possible.

4.8.4   Example intersection and bounding box programs

The following code demonstrates how an intersection and bounding box program combine to describe a simple geometric primitive. The sphere is a simple analytic shape with a well-known ray intersection algorithm. In the following code example, the `sphere` variable encodes the center and radius of a three-dimensional sphere in a `float4`:

Listing 4.32

```
rtDeclareVariable( float4, sphere, , );
rtDeclareVariable( Ray, ray, rtCurrentRay, );
rtDeclareVariable( float3, normal, attribute normal );

RT_PROGRAM void intersect_sphere( int prim_index )
{
    float3 center = make_float3( sphere.x, sphere.y, sphere.z );
    float radius = sphere.w;
    float3 O = ray.origin - center;
    float b = dot( O, ray.direction );
    float c = dot( O, O ) - radius*radius;
    float disc = b*b - c;
    if ( disc > 0.0f ) {
        float sdisc = sqrtf( disc );
        float root1 = (-b - sdisc);
        bool check_second = true;
        if ( rtPotentialIntersection( root1 ) ) {
            normal = (O + root1*D) / radius;
            if ( rtReportIntersection( 0 ) )
                check_second = false;
        }
        if ( check_second ) {
            float root2 = (-b + sdisc);
            if ( rtPotentialIntersection( root2 ) ) {
                normal = (O + root2*D) / radius;
                rtReportIntersection( 0 );
            }
        }
    }
}
```

Note that this intersection program ignores its `prim_index` argument and passes a material index of 0 to `rtReportIntersection`; it represents only the single primitive of its corresponding Geometry object.

The bounding box program for the sphere is very simple:

Listing 4.33

```
RT_PROGRAM void bound_sphere( int, float result[6] )
{
    float3 cen = make_float3( sphere.x, sphere.y, sphere.z );    Define center
    float3 rad = make_float3( sphere.w, sphere.w, sphere.w );    and radius from
                                                                 declared
                                                                 sphere variable

    float3 min = cen - rad;    Compute the minimal and maximal corners of the axis-aligned
    float3 max = cen + rad;    bounding box

    result[0] = min.x;
    result[1] = min.y;
    result[2] = min.z;
    result[3] = max.x;    Store results in order
    result[4] = max.y;
    result[5] = max.z;
}
```

4.9   Selector programs

Ray traversal invokes selector *visit programs* upon encountering a Selector node to programmatically select which of the node's children the ray shall visit. A visit program dispatches the current ray to a particular child by calling the `rtIntersectChild` function. The argument to `rtIntersectChild` selects the child by specifying its index in the range $[0, N)$, where N is given by the argument to `rtSelectorSetChildCount`.

4.9.1   Selector visit program function signature

In CUDA C, visit programs return `void`, take no parameters, and use the `RT_PROGRAM` qualifier:

Listing 4.34

```
RT_PROGRAM void visit_program( void );
```

4.9.2   Example visit program

Visit programs may implement, for example, sophisticated level-of- detail systems or simple selections based on ray direction. The following code sample demonstrates an example visit program that selects between two children based on the direction of the current ray:

Listing 4.35

```
rtDeclareVariable( Ray, ray, rtCurrentRay, );

RT_PROGRAM void visit( void )
{
```

```
    unsigned int index = (unsigned int)( ray.direction.y < 0 );
    rtIntersectChild( index );
}
```

4.10   Callable programs

Callable programs allow for additional programmability within the standard set of OptiX programs. Callable programs are referenced by handles that are set using `RTvariable` or `RTbuffer` on the host. This allows the changing of the target of a function call at run time to achieve, for example, different shading effects in response to user input or customize a more general program based on the scene setup. Also, if you have a function that is invoked from many different places in your OptiX node graph, making it an `RT_CALLABLE_PROGRAM` can reduce code replication and compile time, and potentially improve run time through increased warp utilization.

There are three pieces of callable programs. The first is the program you wish to call. The second is a declaration of a proxy function used to call the callable program. The third is the host code used to associate a callable program with the proxy function that will call it within the OptiX node graph.

Callable programs come in two variants, bound and bindless. Bound programs are invoked by direct use of a program bound to a variable through the host API and inherit the semantic type and variable scope lookup as the calling program. Bindless programs are called via an ID obtained from the `RTprogram` on the host and unlike bound programs do not inherit the semantic type or scope lookup of the calling program

4.10.1   Defining a callable program in CUDA

Defining an `RT_CALLABLE_PROGRAM` is similar to defining an `RT_PROGRAM`:

Listing 4.36

```
RT_CALLABLE_PROGRAM float3 get_color( float3 input_color, float scale)
{
    uint2 tile_size = make_uint2(launch_dim.x/N, launch_dim.y/N);
    if (launch_index.x/tile_size.x ^ launch_index.y/tile_size.y)
        return input_color;
    else
        return input_color * scale;
}
```

`RT_CALLABLE_PROGRAM`s can take arguments and return values just like other functions in CUDA, whereas `RT_PROGRAM`s must return `void`.

4.10.2   Using a callable program variable in CUDA

To invoke an `RT_CALLABLE_PROGRAM` from inside another `RT_PROGRAM`, you must first declare its handle. The handles can be one of two types, `rtCallableProgramId` or `rtCallableProgramX`. Both of these types are templated on the return type followed by the argument types (up to 10 arguments are supported as of OptiX 3.6). The difference between these two will be discussed later in this section.

Listing 4.37

```
typedef rtCallableProgramId<int(int)> callT;
rtDeclareVariable(callT, do_work, ,);
typedef rtCallableProgramX<float(int,int)> call2T;
rtDeclareVariable(call2T, do_more_work, ,);
```

OptiX versions 3.5 and older declared callable programs via the `rtCallableProgram` macro. This macro still works for compatibility, but for SM_20 and newer targets `rtCallableProgram` now creates a declaration similar to `rtCallableProgramX`.

Listing 4.38

```
rtCallableProgram(return_type, function_name, (argument_list) );
```

> **Note:** The third argument must be contained in parentheses.

It is recommended to replace all uses of the macro version of `rtCallableProgram` with the templated version, `rtCallableProgramX`. In addition, if the preprocessor macro RT_USE_TEMPLATED_RTCALLABLEPROGRAM is defined then the old `rtCallableProgram` macro is supplanted by a definition that uses `rtCallableProgramX`.

Listing 4.39

```
#include <optix_world.h>                                Before
rtCallableProgram(int, func, (int,float));
```

```
#define RT_USE_TEMPLATED_RTCALLABLEPROGRAM
#include <optix_world.h>
rtDeclareVariable(                                      After
    rtCallableProgram<int(int,float)>, func, , );
```

Once the program variable is declared, your OptiX program may invoke `function_name` as if it were a standard CUDA function. For example:

Listing 4.40

```
rtDeclareVariable(
    rtCallableProgramId<float3(float3,float)>, get_color,,);

RT_PROGRAM camera()
{
    float3 initial_color, final_color;
    // ... trace a ray, get the initial color ...
    final_color = get_color( initial_color, 0.5f );
    // ... write new final color to output buffer ...
}
```

Because the target of the `get_color` program variable is specified at run time by the host, the `camera` function does not need to take into account how its colors are being modified by the `get_color` function.

In addition to declaring single `rtCallableProgramId` variables, you can also declare a buffer of them, pass them to other functions, and store them for later use.

4.10.3   Setting a callable program on the host

To set up an `RT_CALLABLE_PROGRAM` in your host code, load the PTX function using `rtProgramCreateFromPTXFile`, just like you would any other OptiX program. The resulting `RTprogram` object can be used in one of two ways. You can use the object directly to set an `RTvariable` via `rtVariableSetObject`. This is done for `rtCallableProgramX` and `rtCallableProgram` declared variables.

Alternatively, an ID for the `RTprogram` can be obtained through `rtProgramGetId`. This ID can be used to set the value of a `rtCallableProgramId` typed `RTvariable` (via an `rtVariableSet` function) or the values in a `RTbuffer` declared with type `RT_FORMAT_PROGRAM_ID`. For example:

Listing 4.41

```
RTprogram color_program;
RTvariable color_program_variable;

rtProgramCreateFromPTXFile(
    context, ptx_path, "my_color_program", &color_program );
rtProgramDeclareVariable(
    camera_program, "get_color", &color_program_variable );
```

`rtVariableSetObject(` `    color_program_variable, color_program );`	For `rtCallableProgramX` and `rtCallableProgram`

`int id;`
`rtProgramGetId( color_program, &id );`
`rtVariableSet1i( color_program_variable, id );`
	For `rtCallableProgramId`

`camera_program["get_color"]->setProgramId(` `    color_program );`	For convenience the C++ wrapper has a setProgramId method that gets the ID and sets the variable with it

---

Here is an example of creating a buffer of `rtCallableProgramId` values using the C++ API. This sets up several programs one of which (`times_multiplier`) makes use of a locally defined `RTvariable` called `multiplier` that is unique to each instance of the program.

Listing 4.42

```
Program plus10 =
    context->createProgramFromPTXFile( ptx_path, "plus10" );
Program minus10 =
    context->createProgramFromPTXFile( ptx_path, "minus10" );
Program times_multiplier2 =
```

```
        context->createProgramFromPTXFile( ptx_path, "times_multiplier" );

    times_multiplier2["multiplier"]->setInt(2);

    Program times_multiplier3 =
        context->createProgramFromPTXFile( ptx_path, "times_multiplier" );

    times_multiplier3["multiplier"]->setInt(3);

    Buffer functions =
        context->createBuffer( RT_BUFFER_INPUT, RT_FORMAT_PROGRAM_ID, 5 );

    context["functions"]->set( functions );

    callableProgramId<int(int)>* f_data =                    Here you can use the host defined
        static_cast<callableProgramId<int(int)>*>            type of callableProgramId<> or int
            (functions->map());
    f_data[0] = callableProgramId<int(int)>(plus10->getId());
    f_data[1] = callableProgramId<int(int)>(plus10->getId());
    f_data[2] = callableProgramId<int(int)>(times_multiplier2->getId());
    f_data[3] = callableProgramId<int(int)>(minus10->getId());
    f_data[4] = callableProgramId<int(int)>(times_multiplier3->getId());
    functions->unmap();

    int* f_data_int = static_cast<int*>(functions->map());
    f_data_int[0] = plus10->getId();
    f_data_int[1] = plus10->getId();
    f_data_int[2] = times_multiplier2->getId();
    f_data_int[3] = minus10->getId();
    f_data_int[4] = times_multiplier3->getId();
    functions->unmap();
```

Buffers created using `RT_FORMAT_PROGRAM_ID` can either cast the mapped pointer to a callableProgramId type or to int as seen above.

4.10.4   Bound versus bindless callable programs

Bound callable programs are defined using either the `rtCallableProgramX` templated class or with the backward compatible `rtCallableProgram` macro. Bound programs are referred to as bound because you bind an `RTprogram` directly to an `RTvariable` that is then used to call the program. Binding a program to a variable enables OptiX to extend certain features to the program. Bound programs can be thought of as an extension to the caller, inheriting the semantic type as well as the `RTvariable` lookup scope based on *where the program variable is called from*. For example, if a callable program is called from a closest-hit program then attributes are available to the callable program as well as being able to call functions such as `rtTrace`. Additionally, OptiX will look up identifiers in your callable program in the same scopes as the OptiX programs that invoke it. For example, if invoked from a closest-hit program the lookup scopes will be program, geometry instance, material, then context, where the program scope is the callable program itself instead of the scope of the caller.

Bindless callable programs, on the other hand, inherit neither a program semantic type nor scope. Their scope is always itself (the `RTprogram` object) then the context regardless of *where the program is invoked from*. This is to enable calling these programs from arbitrary locations. Obtaining the ID via `rtProgramGetId` will mark the `RTprogram` as bindless and this `RTprogram` object can no longer be bound to an `RTvariable` (used with `rtCallableProgramX` or `rtCallableProgram`). Bindless programs can only call callable programs, `rtPrintf`, `rtThrow`, and inlineable CUDA functions. Buffer, texture, and variable accesses also work.

Where the callable program variable is attached to the OptiX node graph determines which callable program is invoked when called from another OptiX program. This follows the same variable lookup method that other `Variable` employ. The only difference is that you cannot specify a default initializer.

4.11 Triangle intersection semantics

4.11.1 Primitive index semantic

The primitive index of the `RTgeometrytriangles` type can be queried using the `rtGetPrimitiveIndex` semantic. `rtGetPrimitiveIndex` provides the primitive index similar to what is normally passed to a custom intersection program as an argument. If an primitive-index offset is specified on the geometry (a `Geometry` or `GeometryTriangles` node), `rtGetPrimitiveIndex` reports the primitive index of the geometry (in the range $[0, N)$ for N primitives) plus the offset. This behavior is equal to what is passed to an intersection program.

The `rtGetPrimitiveIndex` semantic is available in any-hit (page 53), closest-hit (page 52) and intersection (page 56) programs.

4.11.2 Hit information semantics

The kind of the hit can be queried using a number of semantics:

```
bool rtIsTriangleHit()

bool rtIsTriangleHitFrontFace()

bool rtIsTriangleHitBackFace()
```

The semantics are available in any-hit (page 53) and closest-hit (page 52) programs.

4.11.3 Built-in intersection program attributes

The built-in triangle intersection program provides a fixed set of attributes:

Listing 4.43

```
rtDeclareVariable(
    float2, barycentrics, attribute rtTriangleBarycentrics, )
```

The two floats describe the barycentric coordinates of the ray hit with the triangle. As for other attributes, the name for the attribute `barycentrics` is user defined, but the semantic name attribute `rtTriangleBarycentrics` must be used. Any other attributes must be computed by the user in a any-hit (page 53) or closest-hit (page 52) program. This means that attributes that usually were computed in the intersection program need to be computed in the any-hit (page 53) and closest-hit (page 52) program.

Further information can be queried using semantics such as the intersected triangle index via the `rtPrimitiveIndex` semantic.

Good practice suggests implementing attributes in the following manner in any-hit (page 53) and closest-hit (page 52) programs:

Listing 4.44

```
rtDeclareVariable(
    float2, barycentrics, attribute rtTriangleBarycentrics, );

rtBuffer<float3> vertex_buffer;
rtBuffer<float3> normal_buffer;
rtBuffer<float2> texcoord_buffer;
rtBuffer<int3>   vindex_buffer;
rtBuffer<int3>   nindex_buffer;
rtBuffer<int3>   tindex_buffer;

struct Attributes {
    float3 geometric_normal;
    float3 shading_normal;
    float2 texcoord;
};

__device__ Attributes attribute_program()
{
    Attributes attribs;
    const unsigned int primIdx = rtGetPrimitiveIndex();

    if (normal_buffer.size() == 0
            || n_idx.x < 0
            || n_idx.y < 0
            || n_idx.z < 0) {
        attribs.shading_normal = attribs.geometric_normal;
    } else {
        int3 n_idx = nindex_buffer[primIdx];
        float3 n0  = normal_buffer[n_idx.x];
        float3 n1  = normal_buffer[n_idx.y];
        float3 n2  = normal_buffer[n_idx.z];
        attribs.shading_normal =
            normalize(n1 * barycentrics.x +
                    n2 * barycentrics.y +
                    n0 * ( 1.0f - barycentrics.x - barycentrics.y));
    }
    // ...
}


RT_PROGRAM void closest_hit()
{
    Attributes attribs = attribute_program();
```

```
    // ...
}
```

5 Motion blur

The previous chapters have described the software structures and functions that OptiX provides as a foundation for implementing a ray-tracing application. A ray generation program (page 47) is responsible for defining pixels in the output image from the result of rays traced into the scene. It is useful to think of this as analogous to a camera. For example, "Example ray generation program" (page 49) presents a simple pinhole camera model.

However, a photographic image is not made instantaneously; it is created by exposing film to light for a finite period of time. Objects moving quickly enough with respect to the shutter duration will appear as streaks in the photograph. This streaking effect is called *motion blur*. To create "photorealistic" images—images that look like photographs—the camera model must also simulate the artifact of motion blur.

The OptiX API as previously described provides two places where motion blur can be implemented:

1. The ray generation program can define a starting time and a duration for a simulated camera shutter, sampling at random times within the shutter duration.

2. The primitive intersection program can define animated primitives by storing multiple positions and interpolating between them, given a random sampling time. However, there is a gap between the time of ray generation and primitive intersection; some parts of scene traversal triggered by `rtTrace` are, for efficiency, not programmable and remain internal to OptiX.

Beginning with OptiX version 5.0, programmers can specify motion data for `Transform` and `Geometry` nodes; OptiX automatically builds `Acceleration` structures that respect this motion.

The `rtTrace` call was also extended in version 5.0 to take an optional `time` argument for the ray. OptiX automatically evaluates `Transform` and `Geometry` motion at this time when traversing the scene. The time value is then available to user programs for intersection and shading.

In the mathematical expressions of this chapter, lowercase letters represent scalars and vectors, uppercase letters represent matrices. A name in the C++ API is written in a `fixed-font` typeface. A product of a scalar with a scalar or vector is represented by dot, as in $a \cdot v$, a vector multiplied by a matrix is represented by $\times$, as in $v \times T$, and matrix multiplication is represented by adjacency, as in SRT for the the multiplication of S by R and then by T.

5.1 Motion in Geometry nodes

The motion blur implementation in OptiX adds functions to define the simulated camera shutter and a variation of the bounding-box program described in Intersection and bounding box programs (page 56).

5.1.1 Defining motion range for Geometry nodes

The interval within which the simulated shutter is open is defined in OptiX by including time as a factor in the ray generation program. A related concept is the time interval during which a `Geometry` node moves. This interval is called the `Geometry` node's *motion range*. These two intervals may not correspond; for example, the motion range may be larger than necessary to allow later adjustments to the shutter interval within the motion range.

The motion range for a `Geometry` node is is defined by `rtGeometrySetMotionRange`.

Listing 5.1

```
RTresult RTAPI rtGeometrySetMotionRange(
    RTgeometry geometry, float timeBegin, float timeEnd );
```

The time range is inclusive, with `timeBegin` $\leq$ `timeEnd`. It defaults to the range $[0.0, 1.0]$ if not set.

The `Geometry` node's motion within the motion range is defined by two or more positions through which the node moves, similar to key frames in traditional animation. These positions are called *motion steps* and are defined by `rtGeometrySetMotionSteps`.

Listing 5.2

```
RTresult RTAPI rtGeometrySetMotionSteps(
    RTgeometry geometry, unsigned int n );
```

If `rtGeometrySetMotionSteps` is not called, or is passed a value of 1, then the `Geometry` node remains static and the time range is ignored during traversal. Note that both here and in `Transform` nodes, motion is specified with a set of keys representing the end points of connected segments. The simplest definition of motion therefore requires two keys.

5.1.2 Bounding boxes for motion blur

Without motion blur, the bounding box program has the following signature.

Listing 5.3

```
RT_PROGRAM void bounding_box_program (
    int prim_index, float result[6]);
```

When motion blur is enabled, the bounding box program adds an integer argument for the *motion index* as the second argument:

Listing 5.4

```
RT_PROGRAM void motion_blur_bounding_box_program (
    int prim_index, int motion_index, float result[6]);
```

The `motion_index` argument is an integer in the range [0, *motion-steps* - 1], where *motion-steps* has been defined by `rtGeometrySetMotionSteps`. The relationship of *motion-steps* to time *t* is:

$$t = \textit{motion-index} \cdot \frac{\textit{time-end} - \textit{time-begin}}{\textit{motion-steps}} \tag{1}$$

A `Geometry` node with more than one motion step must have a bounds program that takes this argument for the motion index. This requirement is enforced during context validation. However, static geometry can use either form of the bounding-box program.

The bounding-box program is responsible for returning the bounding box at the motion index. The set of bounding boxes for a given primitive will be interpolated linearly by OptiX when building and traversing a bounding volume hierarchy. If you want to do something other than linear interpolation later in the intersection program, you must pad the bounding boxes so that when linearly interpolated during traversal, they still bound the nonlinear motion path of the primitive.

Currently only custom primitives can have motion, not "built-in" triangles used optionally for some builders, for example, the `Trbvh` builder.

5.1.3 Border modes

OptiX defines the treatment of `Geometry` nodes evaluated outside its time range using *border modes*:

Listing 5.5

```
RTresult RTAPI rtGeometrySetMotionBorderMode(
    RTgeometry geometry,
    RTmotionbordermode beginMode, RTmotionbordermode endMode );
```

The two border modes can be applied separately for `timeBegin` and `timeEnd`:

`RT_MOTIONBORDERMODE_CLAMP`
> This is the default border mode. The `Geometry` node exists at times less than `timeBegin` or greater than `timeEnd`, with the associated bounding box clamped to its value at `timeBegin` or `timeEnd`, respectively.

`RT_MOTIONBORDERMODE_VANISH`
> The geometry vanishes for times less than `timeBegin` or greater than `timeEnd`.

5.1.4 Acquiring Geometry motion parameter values

The following functions return the values of motion parameters set on `Geometry` nodes:

> `rtGeometryGetMotionSteps`

> `rtGeometryGetMotionBorderMode`

> `rtGeometryGetMotionRange`

5.2 Motion in Acceleration nodes

An `Acceleration` attached to a `GeometryGroup` automatically becomes a motion BVH (or NoAccel) if any of its `Geometry` has more than one motion step. A top level `Acceleration` for

a `Group` becomes a motion BVH if anything in the scene under it (`Transform` or `Geometry` nodes) has motion. Not all types of `Acceleration` support motion; those that do not will throw an exception.

OptiX by default will use two motion steps for a motion BVH even if input `Geometry` or `Transform` nodes have more than two motion steps. Users can change this value via the `motion_steps` property. For example, to use three motion steps in the BVH:

Listing 5.6

```
rtAccelerationSetProperty( accel, "motion_steps", "3" );
```

The value of the `motion_steps` property has several performance implications:

- The value of `motion_steps` must be an integer greater than 0. Setting `motion_steps` to 1 is valid and will produce a static BVH over the union of input bounding boxes.

- Device memory for the BVH scales linearly with `motion_steps`.

- The internal time range for the BVH is the maximum time range over all the child `Geometry` and `Transform` nodes, and cannot be overridden. For performance reasons, it may be better for you to split parts of the scene with different time ranges into different `Acceleration` objects, to minimize empty space in each, but OptiX does not require this.

5.3   Motion in Transform nodes

Motion is added to `Transform` nodes using a set of keys uniformly distributed over a time range. The function `rtTransformSetMotionKeys` defines the key values.

Listing 5.7

```
RTresult RTAPI rtTransformSetMotionKeys(
    RTtransform transform, unsigned int n, RTmotionkeytype type,
    const float* keys );
```

The beginning and ending times within which the motion keys are in effect are defined by the function `rtTransformSetMotionRange`.

Listing 5.8

```
RTresult RTAPI rtTransformSetMotionRange(
    RTtransform transform, float timeBegin, float timeEnd );
```

Motion keys are set by a single assignment and replace any existing data set with previous calls to `rtTransformSetMatrix` or `rtTransformSetMotionKeys`.

5.3.1 Key types

A motion key is defined by either a 3×4 matrix or by a 16-element array that encodes scaling, rotation, and translation.

5.3.1.1 Key type `RT_MOTIONKEYTYPE_MATRIX_FLOAT12`

A `RT_MOTIONKEYTYPE_MATRIX_FLOAT12` key is a 12-float 3×4 matrix in row major order (3 rows, 4 columns). When transforming points, vectors and normals at time t during scene traversal, OptiX will linearly interpolate the two matrices that bracket t to get a matrix M, then apply M, M^{-1}, or their transposes.

5.3.1.2 Key type `RT_MOTIONKEYTYPE_SRT_FLOAT16`

A `RT_MOTIONKEYTYPE_SRT_FLOAT16` key can represent a smooth rotation with fewer keys. Each key is constructed from elements taken from a matrix, a quaternion, and a translation.

A scaling matrix S
> The upper nine elements of an upper triangular 4×4 matrix in row-major order. This matrix can include scale, shear, and a translation. The translation can, for example, define a pivot point for rotation, specified as $p_x p_y p_z$ in matrix S:

$$S = \begin{bmatrix} s_x & a & b & p_x \\ 0 & s_y & c & p_y \\ 0 & 0 & s_z & p_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \tag{2}$$

A quaternion R
> A rotation by an angle θ in quaternion format with angular component $qw = \cos(\theta/2)$ and other components $\begin{bmatrix} q_x & q_y & q_z \end{bmatrix} = \sin(\theta/2) \cdot \begin{bmatrix} a_x & a_y & a_z \end{bmatrix}$ where the axis $\begin{bmatrix} a_x & a_y & a_z \end{bmatrix}$ is normalized.

$$R = \begin{bmatrix} q_x & q_y & q_z & q_w \end{bmatrix} \tag{3}$$

A translation matrix T

A translation defined by t_x, t_y, and t_z components to be applied after the rotation.

$$T = \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \tag{4}$$

The transformations are applied in the order S, then R, then T. For column vectors p and p', $p' = T \times R \times S \times p$.

OptiX will store and interpolate R as a quaternion. Other components are interpolated linearly during traversal. For this type, the dimension of the keys array is $16 \times N$ for N keys. A single key would have these components that refer to the elements of S, R, and T.

$$key = \begin{bmatrix} s_x & a & b & p_x & s_y & c & p_y & s_z & p_z & q_x & q_y & q_z & q_w & t_x & t_y & t_z \end{bmatrix}$$

Note that the order of the elements in this key is based on a serialization of the component values of S (first nine, by row), R, and T.

When transforming points, vectors, and normals at time t, OptiX will effectively first interpolate S, T, and R, then build a combined transform $C = T \times R \times S$ at time t, then apply C, C^{-1}, or their transposes.

5.3.2   Border modes for Transform nodes

As with `Geometry` nodes, you also can define how a `Transform` node is evaluated outside its time range. For example, a time value of 1.2 is passed to `rtTrace` but the `Transform` node has a time range of $[0.0, 1.0]$.

Listing 5.9

```
RTresult RTAPI rtTransformSetMotionBorderMode( RTtransform transform,
    RTmotionbordermode beginMode, RTmotionbordermode endMode );
```

The arguments are the same as for `Geometry`: `RT_MOTIONBORDERMODE_CLAMP` or `RT_MOTIONBORDERMODE_VANISH`. For transforms, `RT_MOTIONBORDERMODE_VANISH` means that the sub-tree under the transform is ignored when the specified time is outside the motion time range.

5.3.3 Acquiring Transform motion parameter values

The following functions return the values of motion parameters set on `Transform` nodes:

> `rtTransformGetMotionBorderMode`
>
> `rtTransformGetMotionKeyCount`
>
> `rtTransformGetMotionKeys`
>
> `rtTransformGetMotionKeyType`
>
> `rtTransformGetMotionRange`

5.4 Examples of motion transforms

The hollowing are examples of the use of motion transforms.

Case 1: Translation by $\begin{bmatrix} t_x & t_y & t_z \end{bmatrix}$
  Use the `MATRIX_FLOAT12` key type with two keys. Set the first key to the identity matrix and the second key to a 3×4 translation matrix with $\begin{bmatrix} t_x & t_y & t_z \end{bmatrix}$ as the rightmost column.

Case 2: Rotation about the origin, with spherical interpolation
  Use the `SRT_FLOAT16` key type with two keys. Set the first key to identity values. For the second key, define a quaternion from an axis and angle, for example, a 60-degree rotation about the z axis is given by:

$$q = \begin{bmatrix} 0 & 0 & \sin(\pi/6) & \cos(\pi/6) \end{bmatrix} \tag{5}$$

Case 3: Rotation about a pivot point, with spherical interpolation of rotation
  Use the `SRT_FLOAT16` key type with two keys. Set the first key to identity values. Represent the pivot as a translation P, and define the second key as follows:

$$S' = P_{inv}S \tag{6}$$
$$T' = TP \tag{7}$$
$$p' = T'RS' \times p \tag{8}$$

Case 4: Scaling about a pivot point

Represent the pivot as a translation $G = \begin{bmatrix} G_x & G_y & G_z \end{bmatrix}$ and modify the pivot point described above:

$$
\begin{aligned}
P'_x &= P_x + (-S_x \cdot G_x + G_x) \\
P'_y &= P_y + (-S_y \cdot G_y + G_y) \\
P'_z &= P_z + (-S_z \cdot G_z + G_z)
\end{aligned}
\tag{9}
$$

5.5   Motion in user programs

A variant of the `rtTrace` function sets the time for the traced ray through an additional argument.

Listing 5.10
| ```
rtTrace( ..., float time = rtCurrentTime );
``` |

Time is not necessarily in the interval of $[0.0, 1.0]$; the interpretation as a relative or absolute time is up to the application. If time is not given, it defaults to one of the following:

- The time of the parent ray that triggered the program; or

- 0.0 if there is no parent ray, for example, in a ray-generation program

In typical use, the sample time is fixed for an entire ray tree. To achieve this, the ray-generation program would specify a sample time for each primary ray traced, but closest-hit programs would not need to specify a time when tracing secondary rays.

The time defined for the current ray can be read from the `rtCurrentTime` semantic variable. Reading `rtCurrentTime` is supported in all programs where `rtCurrentRay` is supported.

The value for the current time is also used by these four functions that set or acquire transform values:

```
rtTransformPoint
```

```
rtTransformNormal
```

```
rtTransformVector
```

```
rtGetTransform
```

6   Post-processing framework

6.1   Overview

Starting with OptiX 5.0 a post-processing framework has been added to OptiX. The post-processing framework allows to post-process images rendered by an OptiX based renderer. The post-processing framework introduces two new types of objects to the OptiX API:

Post-processing stage

> A post-processing stage usually transforms at least one input buffer into an output buffer by post-processing it. The following post-processing stage types are built-in to the current version of OptiX:
>
> - "Deep-learning-based denoiser" (page 79)
> - "Deep-learning-based SSIM predictor" (page 82)
> - "Simple tone mapper" (page 84)

Command lists

> A command list is an ordered list of post-processing stages and OptiX launches

The post-processing framework can be used by instantiating any number of built-in post-processing stages. Those stages and additionally OptiX launches can then be added to a command list. Command lists can later be executed and will execute the launches and all the post-processing stages in the command list, eventually producing one or more output buffers.

Each post-processing stage operates on a number of inputs and produces at least one output. Subsequent post-processing stages can then take the outputs of a post-processing stage which is earlier in the command list and operate on them. Input and outputs are given as OptiX variables of arbitrary types and are managed by the user.

Post-processing stages can be configured using a fixed set of input and output variables with specific names. Note that those variables, even though they are built in, must still be declared in the application like any other OptiX variable.

The available variables are described in the description of the respective post-processing stages. API calls support iteration over the set of predefined variables of a stage as well as direct access to them by name.

A simple command list would start with a launch command which outputs to an `RTbuffer` object. After that a post-processing stage would be added. The variable named `input_buffer` would be set to the `RTbuffer` which holds the output of the launch command. A variable named `output_buffer` would be set to a new `RTbuffer`. When a second post-processing stage would be added after the first one, its `input_buffer` variable would be set to the same `RTbuffer` as the `output_buffer` of the first stage.

6.2   Post-processing stage

6.2.1   Creating a post-processing stage

To create a post-processing stage, use `rtPostProcessingStageCreateBuiltin`.

Input parameters:

`context`

    The OptiX context to which the post-processing stage belongs. The post-processing stage can only be added to a command list belongs to the same context.

`builtin_name`

    The name of the post-processing stage type. See the description of the currently existing built-in post-processing stages for supported names.

Output parameter:

`stage`

    If the call is successful, the created post-processing stage is stored here.

The call returns an error code.

6.2.2   Querying variables

6.2.2.1   Declaring a named variable

Before getting a named variable, you must first declare it using `rtPostProcessingStageDeclareVariable`.

Input parameters:

`stage`

    The post-processing stage to get the variable from.

`name`

    The name of the variable.

Output parameter:

`v`   If the call is successful, a handle to the newly declared variable is returned.

The call returns an error code.

6.2.2.2   Getting a named variable

To get a named variable, use `rtPostProcessingStageQueryVariable`.

Input parameters:

`stage`

    The post-processing stage to get the variable from.

`name`

    The name of the variable.

Output parameter:

`variable`
>   If the call is successful, the variable is stored here.

The call returns an error code.

After obtaining a variable, you can set it to a value that matches the expected type.

6.2.2.3   Iterating over existing variables

To iterate over all existing variables of a post-processing stage, use
`rtPostProcessingStageGetVariableCount` and `rtPostProcessingStageGetVariable`. After
obtaining a variable, you can query the name and the type of the variable.

> **Note:** You can only iterate over variables declared by the application. You cannot use
> these API calls to find out which variables are supported.

6.3   Command lists

6.3.1   Creating a command list

To create a command list, use `rtCommandListCreate`.

Input parameter:

`context`
>   The OptiX context to which the command list belongs. Only post-processing stages
>   belonging to the same context can be added to the command list.

Output parameter:

`list`
>   If the call is successful, the created command list will be stored here

The call returns an error code.

You can create any number of post-processing stages and use them in the same or different
command lists.

6.3.2   Adding a post-processing stage to a command list

To add a post-processing stage to a command list, use the API call
`rtCommandListAppendPostprocessingStage`.

Input parameters:

`list`
>   The command list to append to.

`stage`
>   The post-processing stage to append to the command list.

`launch_width`
>   This is a hint for the width of the launch dimensions to use for this stage.

`launch_height`
>   This is a hint for the height of the launch dimensions to use for this stage.

The call returns an error code.

6.3.3   Appending a launch to a command list

To append a launch to a command list, use `rtCommandListAppendLaunch2D`.

Appending a launch to a command list has the same effect as calling `rtContextLaunch2D` directly except that the launch is executed as part of the command list and at the position in the command list defined by the order in which stages and launches were added to the command list.

Input parameters:

`list`
    The command list to append to.

`entry_point_index`
    The initial entry point into the kernel.

`launch_width`
    Width of the computation grid.

`launch_height`
    Height of the computation grid.

The call returns an error code.

6.3.4   Finalizing a command list

After adding all stages and launches to the command list, finalize the command list using `rtCommandListFinalize`. This call prepares the command list for execution.

> **Note:** After calling finalize it is still possible to set and change the input and output variables. However it is no longer possible to add stages or launches.

6.3.5   Running a command list

To run a command list, use `rtCommandListExecute`. This function can only be called after `rtCommandListFinalize` has been called. The execution operation runs all launches and stages in an order that is compatible with the command list. Later stages and launches can rely on the previous stages to be finished if they use output variables written to by the earlier stages and launches.

Running a command list will first validate the command list. If the setup of the command list is not valid, then execution fails. This can be the case, for example, when necessary variables have not been set or when the type of the variables is not matching the set of expected types. For example, it is not possible to set a variable to a float type if the expected type is a buffer.

After the execution call has returned, the output variables can be read and used for their respective purposes, for example a rendered image can be displayed.

> **Note:** Variable contents must not be changed during the execution of a command list. Doing this will result in undefined behavior. Also note that a command list can be executed any number of times, but only one execution may be active at a time.

See the `optixDenoiser` example shipped with the OptiX SDK. It demonstrates a full post-processing pipeline with a rendering followed by tone mapping and then denoising.

6.4 Built-in post-processing stages

The following sections describe the current built-in post-processing stages.

6.4.1 Deep-learning-based denoiser

Image areas that have not yet fully converged during rendering will often exhibit pixel-scale grainy *noise* due to the insufficient amount of color information gathered by the renderer.

OptiX can estimate the converged image from a partially converged one, a process called *denoising*. Instead of further improving image quality through a large number of path tracing iterations, the denoiser can produce images of acceptable quality with far fewer iterations by post-processing the image.

The OptiX type name for denoising used as an argument to `rtPostProcessingStageCreateBuiltin` is `DLDenoiser`.

The OptiX denoiser comes with a built-in pre-trained model. The model, represented by a binary blob called *training data*, is the result of training the underlying Deep Learning system with a large group of rendered images in different stages of convergence. Since training needs significant computational resources and obtaining a sufficient amount of image pairs can be difficult, a general-purpose model is included with OptiX. This model is suitable for many renderers in practice, but might not always lead to optimal results when applied to images produced by renderers with different noise characteristics compared to those that were present in the original training data.

You can also create a custom model by training the denoiser with your own set of images and use the resulting training data in OptiX, but this process is not part of OptiX itself. To learn how to generate your own training data based on your renderer's images you can attend the course *Rendered Image Denoising using Autoencoders*, which is part of the NVIDIA Deep Learning Institute.

In general, the pixel color space of an image that is used as input for the denoiser should match that of the images it was trained on, although slight variations such as substituting sRGB with a simple gamma curve, should not have a noticeable impact. The images of the training model included with the OptiX distribution were rendered using a gamma value of 2.2.

Using the denoiser is only possible if an additional shared library is available at run time. This shared library is delivered with the OptiX installer and is named `denoiser.dll` on Windows and `denoiser.so` on Linux. The `cudnn` shared library, which is installed with the OptiX SDK, is required at run time. It is only necessary to deliver these shared libraries if your application supports denoising.

> **Note:** Efficient denoising requires high data transfer rates beyond those provided by PCIe 2.0. You should use a PCIe 3.0 slot for the GPU, in particular for GPUs based on the Volta architecture.

6.4.2   Denoiser variables

The denoiser supports the following variables:

`input_buffer`

> A buffer of type `RTbuffer` which contains values of type `float4` representing a noisy image that is to be denoised. The fourth (alpha) channel of the image is not changed by the denoiser. Note that this buffer must contain values between 0 and 1 for each of the three color channels (for example, as the result of tone-mapping) and should be encoded in sRGB or gamma space with a gamma value of 2.2.

> > **Note:**
> >
> > - An image in linear color space can be tone mapped and converted into the correct gamma space, for example, using the `TonemapperSimple` post-processing stage with gamma set to 2.2, before denoising the image.
> >
> > - Applying image blur techniques before denoising an image can have detrimental effects on the quality of the denoised image, especially for very small blur kernels that change the characteristics of the noise. It is therefore recommended to apply image-space techniques only after denoising.

`output_buffer`

> An `RTbuffer` of type `float4`. It must have the same dimensions as the input buffer as it will be used to store the denoised image.

`input_albedo_buffer` *optional*

> The albedo image represents an approximation of the color of the surface of the object, independent of view direction and lighting conditions. In physical terms, the albedo is a single color value approximating the ratio of radiant exitance to the irradiance under uniform lighting. The albedo value can be approximated for simple materials by using the diffuse color of the first hit, or for layered materials by using a weighted sum of the albedo values of the individual BRDFs. For some objects such as perfect mirrors, the quality of the result might be improved by using the albedo value of a subsequent hit instead. The fourth channel of this buffer is ignored, but must have the same type and dimensions as the input buffer. If not declared then denoising will be done without an albedo buffer. It is also possible to disable the albedo buffer by assigning an empty buffer (size 0) to this variable. This is treated as if it wasn't declared in the first place. It is possible to switch albedo on/off by switching between an empty buffer and an albedo buffer between launches.

`input_normal_buffer` *optional*

> This buffer is expected to contain the surface normals of the primary hit in camera space. The camera space is assumed to be right handed such that the camera is looking down the negative z axis, and the up direction is along the y axis. The x axis points to the right. The normal buffer can only be specified if the albedo buffer is present. The fourth channel of this buffer is ignored. It must have the same type and dimensions as the input buffer. If not declared then denoising will be done without a normal buffer. It is also possible to disable the normal buffer by assigning an empty buffer (size 0) to this variable. This is treated as if it wasn't declared in the first place. It is possible to switch normals on/off by switching between an empty buffer and a normal buffer between launches.

`training_data_buffer` *optional*

> A custom training set to be used for denoising. This must be an `RTbuffer` of type `byte`. If this is not set, the built-in training set will be used.

`blend` *optional*

> The `blend` variable is a value of type `float` between 0.0 and 1.0. A value of 0.0 means the fully denoised image is written to the output buffer. A value of 1.0 means that the original image is written to the output buffer. A value in between will produce a blend between original image and denoised image. This can be used for example to reduce the effect of denoising for early iterations and increase it over time. Use this if denoising early iterations produces unacceptable artifacts with your renderer.

`hdr` *optional*

> A `Uint` variable that defines which training data set is used. When `hdr` is set to:
>
> - `zero` (the default), the low-dynamic-range (LDR) training set is used
> - `non-zero`, the high-dynamic-range (HDR) training set is used.
>
> When HDR data are incorrectly denoised with the LDR training data set, color fringing effects are visible in the denoised image.

`denoise_alpha` *optional*

> A `Uint` variable that defines how the alpha channel is denoised. By default the alpha channel of the source image is copied to the result image. If this variable is set to nonzero, the alpha channel is denoised as well instead of copying it from the noisy input image.

6.4.2.1   Defining the denoiser's maximum memory size

As of OptiX version 5.1, the maximum memory size used by the denoiser can be defined by the variable `maxmem`. The denoiser will try to stay below the defined size by splitting the image to be denoised into tiles and denoising them individually. If the defined memory size does not allow denoising, the denoising operation will fail and leave the image in its original state. However, setting a maximum memory size may result in slower denoising performance.

`maxmem` *optional*

> The maximum denoiser memory size in bytes, defined by the integral part of this float variable.

6.4.2.2   HDR input for denoising

OptiX versions greater than 5.0 support *high dynamic range* (HDR) images. For HDR images, single pixels of very high values, or *fireflies*, need to be removed by filtering when using the training set included with OptiX.

A firefly is a statistical outlier, often the result of the contribution of a light path with high contribution that got sampled with low probability. Firefly filtering should happen directly during rendering, for example, when accumulating the contribution of a newly sampled light path to the output buffer.

The HDR denoising process requires the RGB values in the color buffers to be in a range from zero to 10,000. However, denoising very dark images—in which the HDR RGB values are close to zero—may result in incorrect color values.

6.4.2.3 Denoising performance

For HD and 4K images, denoising produced the following memory footprints and timings:

Size	GPU	Memory	Time
1920x1080	GV100	732 MB	19 ms
	P6000	635 MB	49 ms
	P5000	635 MB	70 ms
3840x2160	GV100	2.9 GB	74 ms
	P6000	2.5 GB	193 ms
	P5000	2.5 GB	282 ms

6.4.2.4 Denoiser limitations

In Optix 5.0 and later, the denoiser has the following limitations:

- The denoiser runs under the first GPU found by the system. A different GPU can be selected by calling the functions `cudaSetDevice()`[6].

- There is no CPU fallback for denoising.

- Objects behind transparent surfaces (for example, simulations of glass) will not denoise correctly.

- Denoising produces flickering in a series of images rendered as an animation.

- The denoising training set is not tuned for high-frequency imagery, such as hair or other fine detail, and may produced apparently blurred results.

6.4.3 Deep-learning-based SSIM predictor

OptiX provides a method for estimating how far away a partially converged image is from an imaginary converged image. This estimate might be used, for example, to decide when a Monte Carlo render is finished, or to estimate the time to render completion, or for adaptive sampling, or to determine whether denoising is necessary.

The OptiX SSIM predictor is trained by comparing partially converged images to fully converged images, and measuring their difference using the Structural Similarity Image Metric (SSIM). For more information about SSIM, see Structural similarity[7].

The OptiX type name for SSIM prediction is used as an argument to `rtPostProcessingStageCreateBuiltin` is `DLSSIMPredictor`.

Like the Optix denoiser, this post-processing stage comes with a built-in pre-trained model that has many of the same caveats as the denoiser. Be sure to read and understand the limitations of the DLDenoiser post-processing stage. Be aware that the performance of the SSIM predictor model depends on the noise and color characteristics of the renderer used to product the source image.

---

[6]`http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#device-enumeration`
[7]`https://en.wikipedia.org/wiki/Structural_similarity`

Using the SSIM predictor is only possible if an additional shared library is available at run time. This shared library is delivered with the OptiX installer and is named `sim_predictor.dll` on Windows and `ssim_predictor.so` on Linux. The `cudnn` shared library, which is installed with the OptiX SDK, is required at run time. It is only necessary to deliver these shared libraries if your application supports SSIM prediction.

DLSSIMPredictor is not per-pixel. Instead, it generates a *heat map* output at 1:16 scale relative to the input image. That is, each pixel of this post processing stage corresponds to a 16x16 pixel tile of the input image. This means some compensation may be needed for regions of varying noisiness within a 16x16 pixel region.

DLSSIMPredictor supports the following variables:

`input_buffer`

> A buffer of type `RTbuffer` which contains values of type `float4` representing a possibly noisy image that is to be measured. Because the SSIM predictor does not support HDR, the fourth (alpha) channel of the image is ignored.

> > **Note:** This buffer must contain values between 0 and 1 for each of the three color channels (for example, as the result of tone mapping) and should be encoded in sRGB or gamma space with a gamma value of 2.2.
> >
> > An image in linear color space can be tone mapped and converted into the correct gamma space. For example, before running the SSIM predictor use the `TonemapperSimple` post-processing stage with gamma set to 2.2.

`output_buffer`

> An `RTbuffer` of type `float4`. The output of the SSIM predictor is a *heat map* that corresponds to 16x16 tiles in the input image. Therefore, the output buffer must be 1/16th the size of the input image, rounded up to the nearest integer. For example: If the size of an input image is 256x257, the size of the output buffer must be 16x17.

6.4.3.1   Defining the SSIM predictor's maximum memory size

As of OptiX version 5.1, the maximum memory size used by the SSIM predictor can be defined by the variable `maxmem`. The SSIM predictor will try to stay below the given maximum memory size by splitting the image to be measured into tiles and running SSIM prediction on the tiles individually. If the given maximum memory size does not allow SSIM prediction, the operation will fail and an exception will be thrown. Note that setting a maximum memory size may result in slower SSIM predictor performance and/or slightly different results.

`maxmem` *optional*

> The maximum SSIM predictor memory size in bytes, defined by the integral part of this float variable.

6.4.3.2   SSIM predictor performance

For HD and 4K images, the SSIM predictor produced the following memory footprints and timings:

Size	GPU	Memory	Time (ms)
1920x1080	GV100	441 MB	8.7
	P6000	441 MB	9.2
	P100	444 MB	12.3
	P5000	444 MB	21
3840x2160	GV100	1765 GB	23
	P6000	1765 GB	28
	P100	1777 GB	44
	P5000	1777 GB	80

6.4.3.3 SSIM predictor limitations

The SSIM predictor has the following limitations:

- It runs under the first GPU found by the system. A different GPU can be selected by calling the functions `cudaSetDevice()`[8].

- There is no CPU fallback.

6.4.4 Simple tone mapper

Converting an image containing high-dynamic-range colors to values appropriate for a display device is called *tone mapping*. A tone-mapping post-processing stage in OptiX is defined by using `TonemapperSimple` as an argument to `rtPostProcessingStageCreateBuiltin`. The OptiX tone-mapping stage can also implement gamma correction. Tone-mapping can prepare an image for later use with the denoiser.

The tone mapper predefines the following variables:

`input_buffer`
    This should be an image rendered by a renderer, for example a path tracer. It must be of type `RTbuffer` which contains values of type `float4`.

`output_buffer`
    This must be an `RTbuffer` of type `float4`. It must have the same dimensions as the input buffer as it will be used to store the tone mapped image.

`gamma`
    This is a variable of float type and controls the gamma correction, applied to the first three channels of the output values of the tone mapper. This operation is implemented as `pow(value, 1 / gamma)`. The default value is 1.0, so that no gamma correction is applied.

`exposure`
    This is a variable of float type and acts as a simple multiplier that is applied to the first three channels of the input values. The default value is 1.0.

---

[8]`http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#device-enumeration`

7   Building with OptiX

7.1   Libraries

OptiX comes with several header files and supporting libraries, primarily optix and optixu. On Windows, these libraries are statically linked against the C run-time libraries and are suitable for use in any version of Microsoft Visual Studio, though only the subset of versions listed in the OptiX release notes are tested. If you wish to distribute the OptiX libraries with your application, the VS redistributables are not required by our DLL.

The OptiX libraries are numbered not by release version, but by binary compatibility. Incrementing this number means that a library will not work in place of an earlier version (for example, `optix.2.dll` will not work when an `optix.1.dll` is requested). On Linux, you will find `liboptix.so` which is a soft link to `liboptix.so.1` which is a soft link to `liboptix.so.Z.Y.Z`, the actual library of OptiX version X.Y.Z. `liboptix.so.1` is the binary compatibility number similar to `optix.1.dll`. On MacOS X, `liboptix.X.Y.Z.dylib` is the actual library, and you will also find a soft link named `liboptix.1.dylib` (again, with the 1 indicating the level of binary compatibility), as well as `liboptix.dylib`.

7.2   Header files

There are two principal methods to gain access to the OptiX API. Including `<optix.h>` in host and device code will give access strictly to the C API. Using `<optix_world.h>` in host and device code will provide access to the C and C++ API as well as importing additional helper classes, functions, and types into the OptiX namespace (including wrappers for CUDA's vector types such as float3).

Sample 5 from the SDK provides two identical implementations using both the C (`<optix.h>`) and C++ (`<optixpp_namespace.h>`) API, respectively. Understanding this example should give you a good sense of how the C++ wrappers work.

The `optixu` include directory contains several headers that augment the C API. The namespace versions of the header files (see the list of files below) place all the classes, functions, and types into the `optix` namespace. This allows better integration into systems which would have had conflicts within the global namespace. Backward compatibility is maintained if you include the old headers. It is not recommended to mix the old global namespace versions of the headers with the new `optix` namespace versions of the headers in the same project. Doing so can result in linker errors and type confusion.

`optix_world.h`
    General include file for the C/C++ APIs for host and device code, plus various helper
    classes, functions, and types all wrapped in the `optix` namespace.

`optix.h`
    General include file for the C API for host and device code.

`optixu/optixu_math_namespace.h`

    Provides additional operators for CUDA's vector types as well as additional functions such as `fminf`, `refract`, and an ortho normal basis class.

`optixu/optixpp_namespace.h`

    C++ API for OptiX (backward compatibility with `OptiXu` namespace is provided in `optixpp.h`).

`optixu/optixu_matrix_namespace.h`

    Templated multi-dimensional matrix class with certain operations specialized for specific dimensions.

`optixu/optixu_aabb_namespace.h`

    Axis-aligned bounding box class.

`optixu/optixu_math_stream_namespace.h`

    Standard template library stream operators for CUDA's vector types.

`optixu/optixu_vector_types.h`

    Wrapper around CUDA's `vector_types.h` header that defines the CUDA vector types in the OptiX namespace.

`optixu/optixu_vector_functions.h`

    Wrapper around CUDA's `vector_functions.h` header that defines CUDA's vector functions (for example, `make_float3`) into the `optix` namespace.

7.3   PTX generation

Programs supplied to the OptiX API must be written in PTX. This PTX could be generated from any mechanism, but the most common method is to use the CUDA Toolkit's `nvcc` compiler to generate PTX from CUDA C/C++ code.

When `nvcc` is used, make sure the device code bitness is targeted by using the `-m64` flag. The bitness of all PTX given to the OptiX API must be 64-bit.

When using `nvcc` to generate PTX output specify the `-ptx` flag. Note that any host code in the CUDA file will not be present in the generated PTX file. Your CUDA files should include `<OptiX_world.h>` to gain access to functions and definitions required by OptiX and many useful operations for vector types and ray tracing.

OptiX is not guaranteed to parse all debug information inserted by `nvcc` into PTX files. We recommend avoiding the `--device-debug` flag of `nvcc`. Note that this flag is set by default on debug builds in Visual Studio.

OptiX supports running with NVIDIA Nsight, but does not currently support kernel debugging in Nsight. In addition, it is not recommended to compile PTX code using any `-G` (debug) flags to `nvcc`.

In order to provide better support for compilation of PTX to different SM targets, OptiX uses the `.target` information found in the PTX code to determine compatibility with the currently utilized devices. If you wish your code to run an sm_20 device, compiling the PTX with `-arch sm_30` will generate an error even if no sm_30 features are present in the code. Compiling to sm_20 will run on sm_20 and higher targets.

7.4   SDK build

Our SDK samples' build environment is generated by CMake. CMake is a cross platform tool that generates several types of build systems, such as Visual Studio projects and makefiles. The SDK comes with three text files describing the installation procedures on Windows, Macintosh, and Linux, currently named `INSTALL-WIN.txt`, `INSTALL-MAC.txt` and `INSTALL-LINUX.txt` respectively. See the appropriate file for your operating system for details on how to compile the example source code provided as part of the OptiX SDK.

8 Interoperability with OpenGL

OptiX supports the sharing of data between OpenGL applications and types `rtBuffer` and `rtTextureSampler`. This way, OptiX applications can read data directly from objects such as vertex and pixel buffers, and can also write arbitrary data for direct consumption by graphics shaders. This sharing is referred to as *interoperability* or by the abbreviation *interop*.

8.1 OpenGL interop

OptiX supports interop for OpenGL buffer objects, textures, and render buffers. OpenGL buffer objects can be read and written by OptiX program objects, whereas textures and render buffers can only be read.

8.1.1 Buffer objects

OpenGL buffer objects like PBOs and VBOs can be encapsulated for use in OptiX with `rtBufferCreateFromGLBO`. The resulting buffer is only a reference to the OpenGL data; the size of the OptiX buffer must be set by `rtBufferSetSize1D`, `rtBufferSetSize2D` or `rtBufferSetSize3D`, and the format must be set by `rtBufferSetFormat`. When the OptiX buffer is destroyed, the state of the OpenGL buffer object is unaltered. Once an OptiX buffer is created, the original GL buffer object is immutable, meaning the properties of the GL object like its size cannot be changed while registered with OptiX. However, it is still possible to read and write buffer data to the GL buffer object using the appropriate GL functions. If it is necessary to change properties of an object, first call `rtBufferGLUnregister` before making changes. After the changes are made the object has to be registered again with `rtBufferGLRegister`. This is necessary to allow OptiX to access the objects data again. Registration and unregistration calls are expensive and should be avoided if possible.

8.1.2 Textures and render buffers

OpenGL texture and render buffer objects must be encapsulated for use in OptiX with `rtTextureSamplerCreateFromGLImage`. This call may return `RT_ERROR_MEMORY_ALLOCATION_FAILED` for textures that have a size of 0. Once an OptiX texture sampler is created, the original GL texture is immutable, meaning the properties of the GL texture like its size cannot be changed while registered with OptiX. However, it is still possible to read and write pixel data to the GL texture using the appropriate GL functions. If it is necessary to change properties of a GL texture, first call `rtTextureSamplerGLUnregister` before making changes. After the changes are made the texture has to be registered again with `rtTextureSamplerGLRegister`. This is necessary to allow OptiX to access the textures data again. Registration and unregistration calls are expensive and should be avoided if possible.

Only textures with the following GL targets are supported:

```
GL_TEXTURE_1D

GL_TEXTURE_2D

GL_TEXTURE_2D_RECT
```

```
GL_TEXTURE_3D

GL_TEXTURE_1D_ARRAY

GL_TEXTURE_2D_ARRAY

GL_TEXTURE_CUBE_MAP

GL_TEXTURE_CUBE_MAP_ARRAY
```

Supported attachment points for render buffers are:

```
GL_COLOR_ATTACHMENTn
```

where n is the attachment number.

Not all OpenGL texture formats are supported by OptiX. A table that lists the supported texture formats can be found in the "OpenGL texture formats" (page 119) chapter.

OptiX automatically detects the size, texture format, and number of mipmap levels of a texture.

`rtTextureSamplerSetBuffer` and `rtTextureSamplerGetBuffer` cannot be called for OptiX interop texture samplers and will return `RT_ERROR_INVALID_VALUE`.

9 Interoperability with CUDA

General purpose CUDA programs can be used with OptiX-based ray tracing. For example, you might use a CUDA program before launching OptiX to determine which rays to trace, or to tabulate reflection properties for a material, or to compute geometry. In addition, you may wish to write a CUDA program that post-processes the output of OptiX, especially if OptiX is being used to generate data structures rather than just a rendered image, for example, for computing object or character movement based on visibility and collision rays. These usage scenarios are possible using the OptiX-CUDA interoperability functions described in this chapter.

9.1 Primary CUDA contexts

In order for CUDA and OptiX to interoperate, it is necessary for OptiX and the application to use the same CUDA context. Similar to the CUDA Runtime, OptiX will use the *primary context* for each device, creating it on demand if necessary. Any device pointers that are communicated to and from OptiX will be valid in the primary context. This enables straightforward interoperability of OptiX with both CUDA Runtime API and CUDA Driver API based applications.

Please refer to the CUDA documentation for detailed information about primary contexts.

9.2 Sharing CUDA device pointers

An OptiX buffer internally maintains a CUDA device pointer for each device used by the OptiX context. A buffer device pointer can be retrieved by calling `rtBufferGetDevicePointer`. An application can also provide a device pointer for the buffer to use with `rtBufferSetDevicePointer`. A buffer device pointer can be used by CUDA to update the contents of an OptiX input buffer before launch or to read the contents of an OptiX output buffer after launch. The following example shows how a CUDA kernel can write data to the device pointer retrieved from a buffer:

```
rtBufferGetDevicePointer( buffer, optix_device_ordinal, &device_ptr );
rtDeviceGetAttribute(
    optix_device_ordinal, RT_DEVICE_ATTRIBUTE_CUDA_DEVICE_ORDINAL,
    sizeof(int), &cuda_device_oridinal );
cudaSetDevice( cuda_device_ordinal );
writeData<< blockDim, threadDim >>( device_ptr );
```

Note that each device is assigned an OptiX device ordinal. `rtDeviceGetDeviceCount` can be used to query the number of devices available to OptiX and `rtDeviceGetAttribute` can be used to determine the corresponding CUDA device ordinal for each one (using `RT_DEVICE_ATTRIBUTE_CUDA_DEVICE_ORDINAL`).

9.2.1   Buffer synchronization

Copies of an OptiX buffer's contents may exist on multiple devices and on the host. These copies need to be properly synchronized. For example, if the host copy of a buffer's contents are not up-to-date, a call to `rtBufferMap` may require a copy from a device. If the buffer is an input or input/output buffer, then sometime between the call to `rtBufferUnmap` and `rtContextLaunch` modified host data must be copied to each device used by OptiX. With a multi-GPU OptiX context, getting or setting a buffer pointer for a single device may also require copies to other devices to synchronize buffer data.

9.2.1.1   Automatic single-pointer synchronization

If an application gets or sets a pointer for a single device only, OptiX always assumes that the application has modified the contents of the device pointer and will perform any required synchronizations to other devices automatically. The only exception to this assumption is after a call to `rtBufferUnmap`. If synchronization from the host data to the devices is required, it will override synchronization between devices. Therefore, an application should not modify the contents of a buffer device pointer between a call to `rtBufferUnmap` on the buffer and the next call to `rtContextLaunch`.

9.2.1.2   Manual single-pointer synchronization

If a buffer's contents are not changing for every launch, then the per-launch copies of the automatic synchronization are not necessary. Automatic synchronization can be disabled when creating a buffer by specifying the `RT_BUFFER_COPY_ON_DIRTY` flag. With this flag, an application must call `rtBufferMarkDirty` for synchronizations to take place. Calling `rtBufferMarkDirty` after `rtBufferUnmap` will cause a synchronization from the buffer device pointer at launch and override any pending synchronization from the host.

9.2.1.3   Multi-pointer synchronization

If OptiX is using multiple devices it performs no synchronization when an application retrieves/provides buffer pointers for all the devices. OptiX assumes that the application will manage the synchronization of the contents of a buffer's device pointers.

9.2.2   Restrictions

An application must retrieve or provide device pointers for either one or all of the devices used by a buffer's OptiX context. Getting or setting pointers for any other number of devices is an error. Getting pointers for some devices and setting them for others on the same buffer is not allowed.

Calling `rtBufferMap` or `rtBufferMarkDirty` on a buffer with pointers retrieved/set on all of multiple devices is not allowed. Calling `rtBufferSetDevicePointer` on output or input/output buffers is not allowed.

Setting buffer device pointers for devices which are not used by the buffer's OptiX context is not allowed. An application that needs to copy data to/from a CUDA device that is not part of the OptiX context can do so manually using CUDA, for example, by calling cudaMemcpyPeer or cudaMemcpyPeerAsync.

9.2.3   Zero-copy pointers

With a multi-GPU OptiX context and output or input/output buffers, it is necessary to combine the outputs of each used device. Currently one way OptiX accomplishes this is by using CUDA zero-copy memory. Therefore, `rtBufferGetDevicePointer` may return a pointer to zero-copy memory. Data written to the pointer will automatically be visible to other devices. Zero-copy memory may incur a performance penalty because accesses take place over the PCIe bus.

10 OptiXpp: C++ wrapper for the OptiX C API

OptiXpp wraps each OptiX C API opaque type in a C++ class and provides relevant operations on that type. Most of the OptiXpp class member functions map directly to C API function calls. For example, `VariableObj::getContext` wraps `rtVariableGetContext` and `ContextObj::createBuffer` wraps `rtBufferCreate`.

Some functions perform slightly more complex sequences of C API calls. For example:

Listing 10.1

```
ContextObj::createBuffer(unsigned int type, RTformat format, RTsize width)
```

provides in one call the functionality of:

    rtBufferCreate

    rtBufferSetFormat

    rtBufferSetSize1D

See the OptiX API Reference[9] or header file `optixpp_namespace.h` for a full list of the available OptiXpp functions. The usage of the API is described below.

10.1 OptiXpp objects

The OptiXpp classes consist of a `Handle` class, a class for each API opaque type, and three classes that provide attributes to these objects.

10.1.1 Class `Handle`

All classes are manipulated via the reference-counted `Handle` class. Rather than working with a `ContextObj` directly you would use a `Context` instead, which is simply a typedef for `Handle<ContextObj>`.

In addition to providing reference counting and automatic destruction when the reference count reaches zero, the `Handle` class provides a mechanism to create a handle from a C API opaque type, as follows:

Listing 10.2

```
RTtransform t;
rtTransformCreate( my_context, &t );
Transform Tr = Handle::take( t );
```

The converse of `take` is `get`, which returns the underlying C API opaque type, but does not decrement the reference count within the handle.

---

[9]`http://raytracing-docs.nvidia.com/optix/api/html/index.html`

```
Transform Tr;
// ...
rtTransformDestroy( Tr->get( ) );
```

These functions are typically used when calling C API functions, though such is rarely necessary since OptiXpp provides nearly all OptiX functionality.

10.1.2   Attribute classes

The attributes classes are are `APIObj`, `DestroyableObj`, and `ScopedObj`.

10.1.2.1   Attribute `APIObj`

All object types have the API attribute. This attribute provides the following functions to objects:

`getContext`
> Return the context to which this object belongs.

`checkError`
> Check the given result code and throw an error with appropriate message if the code is not `RT_SUCCESS`. `checkError` is often used as a wrapper around a call to a function that makes OptiX API calls:

Listing 10.4

```
my_context->checkError( sutilDisplayFilePPM( ... ) );
```

10.1.2.2   Attribute `DestroyableObj`

This attribute provides the following functions to objects:

`destroy`
> Equivalent to rt<*type*>Destroy, for example, `rtContextDestroy`.

`validate`
> Equivalent to rt<*type*>Validate, for example, `rtContextValidate`.

10.1.2.3   Attribute `ScopedObj`

This attribute applies only to API objects that are containers for type `RTvariable`. It provides functions for accessing the contained variables. The most basic access is via `operator[]`, as follows:

Listing 10.5

```
my_context["new_variable"]->setFloat( 1.0f );
```

This access returns the variable, but first creates it within the containing object if it does not already exist.

This array operator syntax with the string variable name argument is probably the most powerful feature of OptiXpp, as it greatly reduces the amount of code necessary to access a variable.

The following functions are also available to `ScopedObj` objects:

declareVariable
    Declare a variable associated with this object.

queryVariable
    Query a variable associated with this object by name.

removeVariable
    Remove a variable associated with this object.

getVariableCount
    Query the number of variables associated with this object, typically so as to iterate over them.

getVariable
    Query variable by index, typically while iterating over them.

10.1.2.4 Attributes of OptiXpp objects

The following table lists all of the OptiXpp objects and their attributes.

Object	API	Destroyable	Scoped
Context	✓	✓	✓
Program	✓	✓	✓
Buffer	✓		
Variable	✓		
TextureSampler	✓	✓	
Group	✓	✓	
GeometryGroup	✓	✓	
GeometryInstance	✓	✓	✓
Geometry	✓	✓	✓
Material	✓	✓	✓
Transform	✓	✓	
Selector	✓	✓	

Table 9 – OptiXpp opaque types and their attributes

10.1.3   API objects

In addition to the methods provided by the attribute classes that give commonality to the different API objects each object type also has a unique set of methods. These functions cover the complete set of functionality from the C API, although not all methods will be described here. See the header file `optixpp_namespace.h` for the complete set.

10.1.3.1   Context

The `Context` class provides `create<type>` functions for creating all other opaque types. These are owned by the context and handles to the created object are returned:

Listing 10.6

```
Context my_context;
Buffer Buf =
    my_context->createBuffer(
        RT_BUFFER_INPUT, RT_FORMAT_FLOAT4, 1024, 1024);
```

The `Context` class also provides launch functions, with overloads for 1D, 2D, and 3D kernel launches. It provides many other functions that wrap `rtContext<type>` calls.

10.1.3.2   Buffer

The `Buffer` class provides a map call that returns a pointer to the buffer data, and provides an unmap call. It also provides `set` and `get` functions for the buffer format, element size, and 1D, 2D, and 3D buffer size. Finally, it provides `registerGLBuffer` and `unregisterGLBuffer`.

10.1.3.3   Variable

The `Variable` class provides `getName`, `getAnnotation`, `getType`, and `getSize` functions for returning properties of the variable. It also contains a multitude of `set<type>` functions that set the value of the variable and its type, if the type is not already set:

Listing 10.7

```
my_context["my_dim3"]->setInt( 512, 512, 1024 );
```

The `Variable` object also offers set functions for setting its value to an API object, and provides `setUserData` and `getUserData`.

10.1.3.4   TextureSampler

The `TextureSampler` class provides functions to set and get the attributes of an `RTtexturesampler`, such `setWrapMode`, `setMipLevelCount`, etc.

It also provides `setBuffer`, `getBuffer`, `registerGLTexture`, and `unregisterGLTexture`.

10.1.3.5   Group and GeometryGroup

The remaining API object classes are for OptiX node types. They offer member functions for setting and querying the nodes to which they attach.

The Group class provides `setAcceleration`, `getAcceleration`, `setChildCount`, `getChildCount`, `setChild`, and `getChild`.

10.1.3.6   GeometryInstance

`RTgeometryinstance` is a binding of `Geometry` and `Material`. Thus, `GeometryInstance` provides functions to set and get both the `Geometry` and the `Material`. This includes `addMaterial`, which increments the material count and appends the given `Material` to the list.

10.1.3.7   Geometry

Six unique functions provided by the `Geometry` class define the `set` and `get` methods for the components of the class:

```
setBoundingBoxProgram

getBoundingBoxProgram

setIntersectionProgram

getIntersectionProgram

setPrimitiveCount

getPrimitiveCount
```

The `Geometry` class also provides functions `markDirty` and `isDirty`.

10.1.3.8   Material

A `Material` consists of a closest-hit (page 52) and an any-hit (page 53) program, and is a container for the variables appertaining to these programs. It contains set and get functions for these programs:

```
setClosestHitProgram

getClosestHitProgram

setAnyHitProgram

getAnyHitProgram
```

10.1.3.9   Transform

An `RTtransform` node applies a transformation matrix to its child, so the `Transform` class offers `setChild`, `getChild`, `setMatrix`, and `getMatrix` methods.

10.1.3.10   Selector

A `Selector` node applies a visit (page 59) program to operate on its multiple children. The `Selector` class includes the following functions to set and get components of the class:

```
setVisitProgram

getVisitProgram

setChildCount
```

```
getChildCount

setChild

getChild
```

10.1.4   Exception

The `Exception` class of OptiXpp encapsulates an error message. These errors are often the direct result of a failed OptiX C API function call and subsequent `rtContextGetErrorString` call. Nearly all methods of all object types can throw an exception using the `Exception` class. Likewise, the checkError function can throw an `Exception`.

Additionally, the `Exception` class can be used explicitly by user code as a convenient way to throw exceptions of the same type as OptiXpp.

Call `makeException` to create an Exception.

Call `getErrorString` to return a `std::string`[10] for the error message as returned by `rtContextGetErrorString`.

---

[10]`https://en.cppreference.com/w/cpp/string/basic_string`

11   OptiX Prime: low-level ray tracing API

11.1   Overview

OptiX is generally used to represent an entire algorithm implementation, whether that be rendering, visibility, radiation transfer, or anything else. The many user programmable portions of OptiX allow the application to express complex operations, such as shading, that are tightly intermingled, often recursively, with the ray tracing operations and expressed in a single-ray programming model. By encapsulating the programmable portions of the algorithm and owning the entire algorithm, OptiX can execute the entire algorithm on the GPU and optimize the execution for each new GPU as it is released.

Sometimes the algorithm as a whole does not benefit from this tight coupling of user code and ray tracing code, and only the ray tracing functionality is needed. Visibility, trivial ray casting rendering, and ray tracing very large batches of rays in phases may have this property. OptiX Prime is a set of OptiX APIs designed for these use cases. Prime is specialized to deliver high performance for intersecting a set of rays against a set of triangles. Prime is a thinner, simpler API, since programmable operations, such as shading, are excluded. Prime is also suitable for some quick experimentation and hobby projects.

The OptiX Prime API consist of four main objects:

`BufferDesc`
> Wraps application managed buffers and provides descriptive information about them.

`Context`
> Manages resource allocation.

`Model`
> Represents a set of triangles and an acceleration structure.

`Query`
> Coordinates the intersection of rays with a model.

11.2   Context

An OptiX Prime context performs two main functions. The first function is to manage objects created by the API. The context can create objects, some of which can also create other objects. All of these objects are registered with the context and will be destroyed when the context is destroyed. The second function is to encapsulate a particular back end that performs the actual computation.

Currently the following context types are supported:

> `RTP_CONTEXT_TYPE_CPU`

> `RTP_CONTEXT_TYPE_CUDA`

`RTP_CONTEXT_TYPE_CPU` is intended to be used as a fallback when an acceptable CUDA device is not available. It will allow an application to run, despite the absence of CUDA-capable

GPUs, but will have lower performance. In certain situations it might also make sense to use CPU and CUDA contexts in parallel.

`RTP_CONTEXT_TYPE_CUDA` by default will only use the fastest available device. It is also possible to specify a specific device (or a list of devices) to be used by the context by supplying a list of device numbers to `rtpContextSetCudaDeviceNumbers`. The fastest device in this list is used as the primary device. Acceleration structures will be built on that primary device and copied to the others. Specified devices must feature compute capability SM 3.0 or greater. All devices will be used for ray tracing with work being distributed proportionally to each device's computational power. Note that this distribution can be rather costly if the rays are stored in device memory though. For maximum efficiency it is recommended to only ever select one device per context. The CUDA context of the primary device is made current after a call to `rtpContextCreate` or `rtpContextSetCudaDeviceNumbers`.

The following code demonstrates how to create a context and specify which devices to use. In this example, a CPU context is created as a fallback.

Listing 11.1

```
RTPcontext context;
if (rtpContextCreate( RTP_CONTEXT_TYPE_CUDA, &context ) == RTP_SUCCESS ) {
    int deviceNumbers[] = {0,1};
    rtpContextSetCudaDeviceNumbers( 2, deviceNumbers );
}
else {
    rtpContextCreate( RTP_CONTEXT_TYPE_CPU, &context );
}
```

11.3   Buffer descriptor

The buffers used to send and receive data from OptiX Prime are managed by the application. A buffer descriptor is an object that provides information about a buffer, such as its format and location, as well as the pointer to the buffer's data. OptiX Prime supports the following buffer types:

> `RTP_BUFFER_TYPE_HOST`
>
> `RTP_BUFFER_TYPE_CUDA_LINEARb`

A buffer descriptor is created by calling `rtpBufferDescCreate`. A buffer of type `RTP_BUFFER_TYPE_CUDA_LINEAR` is assumed to reside on the current CUDA device. The device number for the buffer can be specified explicitly by calling `rtpBufferDescSetCudaDeviceNumber`.

The portion of the buffer to use for input or output is specified by calling `rtpBufferDescSetRange`. The range is specified in terms of the number of elements.

For buffers containing vertex data, it is possible to specify a stride in bytes between each element. This is useful for vertex buffers that contain interleaved vertex attributes, as shown in the following example:

> *Listing 11.2*
>
> ```
> struct Vertex {
>     float3 pos, normal, color;
> };
> ...
> RTPbufferdesc vertsBD;
> rtpBufferDescCreate(
>     context, RTP_BUFFER_FORMAT_VERTEX_FLOAT3, RTP_BUFFER_TYPE_HOST,
>     verts, &vertsBD);
> rtpBufferDescSetRange(vertsBD, 0, numVerts);
> rtpBufferDescSetStride(vertsBD, sizeof(Vertex));
> ```

The functions that accept buffer descriptors as parameters have *copy semantics*. This means, for example, that when a buffer descriptor is used in a call to set the rays to be queried that also results in a change to the buffer descriptor, that change to the buffer descriptor will not be visible to the query. Changing the contents of the buffer itself could, however, affect the query.

Buffer descriptors are lightweight objects and can be created and destroyed as needed.

Buffers of any type can be passed to OptiX Prime functions and the will automatically be copied to the appropriate location. For example, a vertex buffer on the host will be copied to the primary device when the context is of type `RTP_CONTEXT_TYPE_CUDA`. While convenient, this automatic copying may require the allocation of memory on the device and can negatively impact performance. For best performance it is recommended that the following buffer and context types be used together:

Context type	Buffer type
RTP_CONTEXT_TYPE_CPU	RTP_BUFFER_TYPE_HOST
RTP_CONTEXT_TYPE_CUDA	RTP_BUFFER_TYPE_CUDA_LINEAR (on the primary device)

11.4   Model

A model represents either a set of triangles or a group of model instances, in addition to an acceleration structure built over the triangles or instances. A model is created with an associated context by calling `rtpModelCreate`, and can be destroyed with `rtpModelDestroy`.

11.4.1   Triangle models

Triangle data for the model is supplied by calling `rtpModelSetTriangles` with a vertex buffer descriptor and an optional index buffer descriptor. If no index buffer is supplied then the vertex buffer is considered to be a flat list of triangle vertices, with every set of three vertices forming a triangle (also know as a "triangle soup").

`rtpModelUpdate` creates the acceleration structure over the triangles. It is important that the vertex and index buffers specified in `rtpModelSetTriangles` remain valid until `rtpModelUpdate` is finished. If the flag `RTP_MODEL_HINT_ASYNC` is specified, some or all of the acceleration structure update may run asynchronously and `rtpModelUpdate` may return before the update is finished. `rtpModelFinish` blocks the calling thread until the update is finished. `rtpModelGetFinished` can be used to poll until the update is finished. Once the update has finished, the input buffers can be modified.

The following code demonstrates how to create a model from a vertex buffer with an asynchronous update. The code assumes that a vertex buffer descriptor vertsBD already exists:

Listing 11.3

```
RTPmodel model;
rtpModelCreate(context, &model);
rtpModelSetTriangles(model, 0, vertsBD);
rtpModelUpdate(model, RTP_MODEL_HINT_ASYNC);

// ... Do useful work on CPU while GPU is busy

rtpModelFinish(model);
// Now safe to modify vertex buffer
```

For some use cases, the user may wish to have explicit control over multi-GPU computation rather than using the automatic multi-GPU support provided by OptiX Prime. A context can be created for each device and work can be distributed manually to each context. OptiX Prime provides `rtpModelCopy` to copy a model from one context to another so that it is not necessary to create and update the model in each context. `rtpModelCopy` can also be used to build multiple models in parallel on different devices, then broadcast the results to each device. When using older devices, `rtpModelCopy` can be used to build an acceleration structure in a CPU context and copy it to the context that uses the devices.

Beyond the memory used by the final acceleration structure, some additional memory is needed during `rtpModelUpdate`. The amount used may be controlled by calling `rtpModelSetBuilderParameter`. The `RTP_BUILDER_PARAM_CHUNK_SIZE` controls the amount of scratch space used while building. The minimum scratch space is currently 64MB, and the default scratch space is 10video memory for CUDA contexts, and 512MB for CPU contexts. A chunk size of -1 signifies unlimited. In this case about 152 bytes per triangle are used while building the acceleration structure.

`RTP_BUILDER_PARAM_USE_CALLER_TRIANGLES` controls whether to create a possibly transformed copy of the vertex buffer data, or to use the buffer supplied by the user, thus saving memory. If a model is copied, and the source model is using the user supplied triangle data to save memory, the user triangles will be automatically copied as well. If this is not intended, it is necessary to set `RTP_BUILDER_PARAM_USE_CALLER_TRIANGLES` on the destination model as well, before the copy is performed. Afterward, `rtpModelSetTriangles` must be called to supply the user triangles on the destination model.

11.4.2   Instancing

Using instancing, it is possible to compose complex scenes using existing triangle models. (See "Triangle models" (page 103).) Instancing data for a model is supplied by calling `rtpModelSetInstances` with an instance buffer descriptor and a transformation buffer descriptor. The ranges for these buffer descriptors must be identical. The type of the instance buffer descriptor must be `RTP_BUFFER_TYPE_HOST`, and the format `RTP_BUFFER_FORMAT_INSTANCE_MODEL`. For transformations, the buffer descriptor format can be either `RTP_BUFFER_FORMAT_TRANSFORM_FLOAT4x4` or `RTP_BUFFER_FORMAT_TRANSFORM_FLOAT4x3`. If a stride is specified for the transformations, it must be a multiple of 16 bytes. Furthermore, the matrices must be stored in row-major order.

Please note that only affine transformations are supported, and that the last row is always assumed to be [0.0, 0.0, 0.0, 1.0].

In contrast to triangle models, instance models can not be copied.

The `primeInstancing` example that ships with the OptiX SDK demonstrates the details of instancing a model in OptiX.

11.4.3 Masking

With masking, it is possible to specify per triangle visibility information in combination with a per ray mask. In order to use masking, triangle data must be specified with the `RTP_BUFFER_FORMAT_INDICES_INT3_MASK_INT` buffer format. Furthermore the user-triangle build parameter must be set, as well as ray format `RTP_BUFFER_FORMAT_RAY_ORIGIN_MASK_DIRECTION_TMAX` must be used. The per triangle visibility flags are evaluated by a bitwise AND operation with the currently processed ray's `MASK` field before a ray-triangle intersection is performed. If the result is non-zero, the ray-triangle intersection test is skipped. The combination of a per-triangle mask with a per-ray mask makes it possible to exclude triangles based on different ray generations.

If the per triangle mask values need to be updated, `rtpModelSetTriangles` must be called again, with a successive call to `rtpModelUpdate`. Using the `RTP_MODEL_HINT_MASK_UPDATE` flag indicates that only the per triangle mask has changed, but that no rebuild of the acceleration structure is needed.

For an example of masking, please refer to the primeMasking sample which ships with the OptiX SDK.

11.5 Query

A query is used to perform the actual ray tracing against a model. The query is created from a model using `rtpQueryCreate`. The following types of queries are supported:

```
RTP_QUERY_TYPE_ANY

RTP_QUERY_TYPE_CLOSEST
```

Along any given ray there may be a number of intersection points. `RTP_QUERY_TYPE_CLOSEST` returns the first hit along the ray. `RTP_QUERY_TYPE_ANY` returns the first hit found, whether it is the closest or not. The query takes a buffer of rays to intersect and a buffer to store the resulting hits. There are several formats for the rays and hits. The main advantage of the different formats is that some require less storage than others. This is important for minimizing the transfer time of rays and hit data between the host and the device and between devices.

Once the ray and hit buffers have been specified, the query can be executed by calling `rtpQueryExecute`. The ray buffer must not be modified until after this function returns. If the flag `RTP_QUERY_HINT_ASYNC` is specified, `rtpQueryExecute` may return before the query is actually finished. `rtpQueryFinish` can be called to block the current thread until the query is finished, or `rtpQueryGetFinished` can be used to poll until the query is finished. At this point all of the hits are guaranteed to have been returned, and it is safe to modify the ray buffer.

The following code demonstrates how to execute a query using ray and hit buffers:

Listing 11.4

```
RTPquery query;
RTPbufferdesc raysBD, hitsBD;

rtpQueryCreate(model, RTP_QUERY_TYPE_CLOSEST, &query);   Fill in raysBD with rays
rtpQuerySetRays(query, raysBD);
rtpQuerySetHits(hits, hitsBD);
rtpQueryExecute(query, 0);

···   Now safe to modify ray buffer and process hits
```

With `rtpQuerySetCudaStream`, it is possible to specify a specific CUDA stream which can be used to synchronize (asynchronous) queries and user CUDA kernel launches. If no stream is specified, the CUDA default stream is used.

A query may be executed multiple times. Note that `rtpQueryFinish` and `rtpQueryGetFinished` only apply to the stream corresponding to the last call to `rtpQueryExecute`. Therefore, if the stream has been changed between asynchronous calls to `rtpQueryExecute` it may be necessary to manually synchronize the streams, for example, by calling cudaStreamSynchronize or using CUDA events (see the *CUDA Programming Guide*[11]).

11.6   Utility functions

In addition to the basic objects and their functions, OptiX Prime has several utility functions.

11.6.1   Page-locked buffers

The performance of transfers between devices and the host can be improved by page-locking the host memory. Functions for page-locking already allocated memory are provided in the CUDA Runtime API. For convenience, OptiX Prime provides the functions `rtpHostBufferLock` and `rtpHostBufferUnlock` so that it is possible to achieve better performance with host buffers without having to invoke CUDA functions directly. Note that page-locking excessive amounts of memory may degrade system performance, since it reduces the amount of memory available to the system for paging. As a result, this function is best used sparingly to register staging areas for data exchange between host and device.

11.6.2   Error reporting

All functions in OptiX Prime return an error code. The function `rtpGetErrorString` translates an error code into a string. `rtpContextGetLastErrorString` returns an error string for the last error encountered. This error string may contain additional information beyond a simple error code. Note that this function may also return errors from previous asynchronous launches, or from other threads.

11.7   Multi-threading

The OptiX Prime API is thread safe. It is possible to share a single context among multiple host threads. However only one thread may access a context (or objects associated with it) at a time. Therefore to avoid locking other threads out of the API for extensive periods of time, the asynchronous APIs should be used. Care must also be taken to synchronize state changes to

---

[11]https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html

API objects. For example, if two threads try to set the ray buffer on the same query at the same time, a race condition can occur.

11.8   Streams

By default, all computation in OptiX Prime (for example, updating models and executing queries) takes place within the default CUDA stream of the primary device. However, with `rtpQuerySetCudaStream` it is possible to specify a specific CUDA stream which can be used to synchronize (asynchronous) queries and user CUDA kernel launches.

12 OptiX Prime++: C++ wrapper for the OptiX Prime API

OptiX Prime++ wraps each OptiX Prime C API opaque type in a C++ class and provides relevant operations on that type. Most of the OptiX Prime++ class member functions map directly to C API function calls. For example, `setCudaDeviceNumbers` maps directly to `rtpContextSetCudaDeviceNumbers`.

Some functions perform slightly more complex sequences of C API calls. For example

    setTriangles

provides in one call the functionality of

    rtpBufferDescCreate

    rtpBufferDescSetRange

    rtpBufferDescSetStride

    rtpModelSetTriangles

    rtpBufferDescDestroy

See the OptiX API Reference[12] or header file `optix_primepp.h` for a full list of the available OptiX Prime++ functions. Using the API is described below.

12.1 OptiX Prime++ objects

Manipulation of OptiX Prime objects is performed via reference counted `Handle` classes which encapsulate all OptiX Prime objects functionalities.

All classes are manipulated via the reference-counted `Handle` class. Rather than working with a `ContextObj` directly you would use a `Context` instead, which is simply a typedef for `Handle<ContextObj>`.

12.1.1 Context object

The `Context` object wraps the OptiX Prime C API `RTPcontext` opaque type and its associated function set representing an OptiX Prime context. Its constructor requires a `RTPcontexttype` type in order to specify the type of the context to be created (CPU or GPU). By default, only the fastest GPU will be used. A different device (or a list of devices) can be selected by using `setCudaDeviceNumbers`. Note that for maximum efficiency it is recommended to only ever select one device per context.

---

[12]`http://raytracing-docs.nvidia.com/optix/api/html/index.html`

Listing 12.1

```
Context context( Context::create(contextType) );
```

For all objects the following pattern is also possible:

Listing 12.2

```
Context context;
context = Context::create(contextType);
```

The `Context` also provides functions to create OptiX Prime objects directly owned by it:

> createBufferDesc

> ContextObj::createModel

12.1.2   `BufferDesc` object

The `BufferDesc` object wraps a `RTPbufferdesc` object opaque type and its associated function set representing an OptiX Prime buffer descriptor.

The creation of a `BufferDesc` object is demanded to an owning `Context` object since each `BufferDesc` object needs to be assigned to an OptiX Prime context.

12.1.3   Model object

The `Model` object wraps a `RTPmodel` object opaque type and its associated function set representing an OptiX Prime model.

The creation of a `Model` object is demanded to an owning Context object since each `Model` object needs to be assigned to an OptiX Prime context.

Variants of the `setTriangles` functions are provided to allow creating a model by either a custom-format set of triangles, with a supplied indices buffer descriptor or directly from a supplied vertices buffer descriptor.

12.1.4   Query object

The `Query` object wraps a `RTPquery` object opaque type and its associated function set representing an OptiX Prime query.

The creation of a `Query` object is demanded to an owning `Context` object since each `Query` object needs to be assigned to an OptiX Prime context.

Variants of the `setRays` and `setHits` functions are provided to allow setting the rays and the hits for a query from either a custom-format user supplied buffer or from a buffer descriptor.

12.1.5   Exception class

The `Exception` class provides methods to deal with OptiX Prime exceptions. Both error code and error description methods are provided as well as a wrapper for the `rtpContextGetLastErrorString` function

---

Listing 12.3

```
catch ( Exception& e ) {
std::cerr « "An error occurred with error code "
    « e.getErrorCode() « " and message "
    « e.getErrorString() « std::endl;
}
```

13 Performance guidelines

Subtle changes in your code can dramatically alter performance. This list of performance tips should help when using OptiX.

Where possible use floats instead of doubles.
> This also extends to the use of literals and math functions. For example, use `0.5f` instead of `0.5` and `sinf` instead of `sin` to prevent automatic type promotion. To check for automatic type promotion, search the PTX files for the `.f64` instruction modifier.

Match the launch dimensionality to the problem.
> OptiX will try to partition thread launches into tiles with the same dimensionality as the launch. To have maximal coherency between the threads of a tile you should choose a launch dimensionality that is the same as the coherence dimensionality of your problem. For example, the common problem of rendering an image has 2D coherency (adjacent pixels both horizontally and vertically look at the same part of the scene), so a 2D launch is appropriate. Conversely, a collision detection problem with many agents each looking in many directions may appear to be 2D (the agents in one dimension and the ray directions in another), but there is rarely coherence between different agents, so the coherence dimensionality is one, and performance will be better by using a 1D launch.

Do not build an articulate scene graph with `Group`, `Transform` and `GeometryInstance`.
> Try to make the topology as shallow and minimal as possible. For example, for static scenes the fastest performance is achieved by having a single `GeometryGroup`, where transforms are flattened to the geometry. For scenes where `Transform` are changing all the static geometry should go in one `GeometryGroup` and each `Transform` should have a single `GeometryGroup`. Also, if possible, combine multiple meshes into a single mesh.

Reuse the same program, customized by different variable values.
> Each new `Program` object can introduce execution divergence. Try to reuse the same program with different variable values. However, don't take this idea too far and attempt to create a monolithic program that covers all cases—a so-called "über shader." This will create execution divergence within the program. Experiment with your scene to find the right balance.

Try to minimize live state across calls to `rtTrace` in programs.
> For example, in a closest hit program temporary values *used* after a recursive call to `rtTrace` should be *computed* after the call to `rtTrace`, rather than before, since these values must be saved and restored when calling `rtTrace`, impacting performance. `RTvariable` declared outside of the program body are exempt from this rule.

Choose types that optimize writing to buffers.
> In multi-GPU environments `INPUT_OUTPUT` and `OUTPUT` buffers are stored on the host. In order to optimize writes to these buffers, types of either 4 bytes or 16 bytes (for example, `float`, `uint`, or `float4`) should be used when possible. One might be tempted to make an output buffer used for the screen `float3` for an RGB image. However, using a `float4` buffer instead will result in improved performance.

For example, the `float4` value can be created at the time of assignment to the output buffer:

```
output_buffer[launch_index] = make_float4(result_color);
```

This suggestion also holds for user-defined types. See the `progressivePhotonMap` sample for an example of accessing user-defined structs with `float4s`.

Align memory accesses to structs containing four-vectors.

Memory accesses to structs containing four-vectors, such as `float4`, need to be 16-byte aligned for optimal performance. Make the alignment by placing the largest aligned variables first in structs.

Use a separate buffer copy per device in a multi-GPU environment.

In multi-GPU environments, `INPUT_OUTPUT` buffers may be stored on the device, with a separate copy per device by using the `RT_BUFFER_GPU_LOCAL` buffer attribute. This is useful for avoiding the slower reads and writes by the device to host memory. `RT_BUFFER_GPU_LOCAL` is useful for scratch buffers, such as random number seed buffers and variance buffers.

Use iteration instead of recursion where possible.

Recursion may not be necessary in specific cases of a general algorithm that typically requires it, for example, a path tracing implementation that does not allow ray branching. See the path_tracer sample for an example of how to use iteration instead of recursion when tracing secondary rays.

Use the `rtTransform*` functions.

These functions provides the best performance, in contrast to explicitly transforming the matrix returned by `rtGetTransform`.

Disable exceptions that are not needed.

While it is recommended to turn on all available exception types during development and for debugging, the error checking involved in some operations, for example, to validate buffer index bounds, is usually not necessary in the final product.

Avoid recompiling the OptiX kernel.

Recompilation can be triggered when certain changes to the input programs or variables occur. For example, swapping the `ClosestHit` program of a `Material` between two programs will cause a recompile on each swap because the kernel consists of different code, whereas creating two `Materials`, one with each program, and swapping between the two `Materials` will not cause a recompile because only the node graph is changing, not the code. Creating dummy nodes with the alternate programs is one way to provide all of the code at once. In addition, avoid changing the layout of variables attached to scope objects.

Define a variable *once*.

It is possible for a program to find multiple definitions for a variable in its scopes depending upon where the program is called. Variables with definitions in multiple scopes are said to be *dynamic variables* and may incur a performance penalty.

Initialize variables.

Uninitialized variables can increase register pressure and negatively impact performance.

Use the `--use-fast-math` compile option of `nvcc`.

When creating PTX code using `nvcc`, adding `--use-fast-math` as a compile option can reduce code size and increase the performance for most OptiX programs. This can come at the price of slightly decreased numerical floating point accuracy. See the nvcc documentation[13] for more details.

13.1   Guidelines for OptiX Prime

The following performance guidelines apply to OptiX Prime:

Use page-locked host memory for `RTP_CONTEXT_TYPE_CUDA`.

Use page-locked host memory for host buffers to improve performance when using contexts of type `RTP_CONTEXT_TYPE_CUDA`. Asynchronous API calls involving host buffers can may not actually be asynchronous if the host memory is not page-locked.

Manage multiple GPUs manually to improve PCIe performance.

Multi-GPU contexts, while convenient, are limited by PCIe bandwidth. Ray and hit buffers reside in a single location (either the host or a device) and must be copied over the PCIe bus to multiple devices. It is possible to obtain better performance by managing multiple GPUs manually. By allocating a context for each device and generating rays and consuming hits on the device, transfers can be avoided.

Use a separate context with its own thread for each device.

For maximum concurrency with model updates, use a separate context for each device, each running in its own host thread. There are currently some limitations on the amount of concurrency that can be achieved from a single host thread which will be addressed in a future release.

Initialize contexts when no long-running kernels are active.

Prime contexts allocate a small amount of page-locked host memory. Because the allocation of page-locked memory can sometimes block when other kernels are running, it is best to initialize contexts when no long-running kernels are active.

Create queries from a multi-GPU context with buffers residing in page-locked memory.

With the current implementation, the performance of queries created from a multi-GPU context are generally better when used with buffers residing in page-locked host memory rather than on a device.

Make multi-threaded execution asynchronous.

Use the asynchronous API when using Prime in a multi-threaded setting. The asynchronous API will achieve better concurrency in a multi-threaded setting.

---

[13]`http://docs.nvidia.com/cuda/cuda-compiler-driver-nvcc/index.html`

14 Caveats

Keep in mind the following caveats when using OptiX:

Stack size

Setting a large stack size will consume GPU device memory. Try to minimize the stack as much as possible. Start with a small stack and with the use of an exception program that will make it obvious you have exceeded your memory, increase the stack size until the stack is sufficiently large.

Disallowed features

The use of `__shared__` memory, PTX `bar` or CUDA `syncthreads` within a program is not allowed.

Thread indices

`threadIdx` in CUDA can map to multiple launch indices (for example, to pixels). Use the rtLaunchIndex semantic instead.

Memory and print functions

Use of the CUDA `malloc` and `free` functions within a program is not supported. Attempts to use these functions will result in an illegal symbol error.

Thread safety

Currently, the OptiX host API is *not* guaranteed to be thread-safe. While it may be successful in some applications to use OptiX contexts in different host threads, it may fail in others. OptiX should therefore only be used from within a single host thread.

Fermi texture limitations

On Fermi architecture the number of cube, layered, and half float textures is limited to 127. Mip levels of textures over the 127 limit are stripped to one level.

OptiX 6.0 Programming Guide © 2019 NVIDIA Corporation

15 OpenGL texture formats

The following symbols represent the texture formats that are supported by OptiX to provide interoperability with OpenGL. Some symbols include a suffix that defines data size and other attributes. This suffix is represented by *<type>* in the table. For example, two formats to represent luminance are `GL_LUMINANCE8` and `GL_LUMINANCE32I_EXT`.

Format	*<type>*
`GL_RED`	
`GL_RG`	
`GL_RGBA`	
`GL_LUMINANCE`	
`GL_ALPHA`	
`GL_LUMINANCE_ALPHA`	
`GL_INTENSITY`	
`GL_R`*<type>* `GL_RG`*<type>* `GL_RGBA`*<type>*	8 16 16F 32F 8UI 16UI 32UI 8I 16I 32I
`GL_LUMINANCE`*<type>* `GL_ALPHA`*<type>* `GL_LUMINANCE_ALPHA`*<type>* `GL_INTENSITY`*<type>*	8 16 16F_ARB 32F_ARB 8UI_EXT 16UI_EXT 32UI_EXT 8I_EXT 16I_EXT 32I_EXT